



HLD Concepts Brief Overview

1. [Vertical vs Horizontal scaling](#)
2. [CAP theorem](#)
3. [ACID vs BASE](#)
4. [C in ACID vs C in CAP](#)
5. [Fault Tolerance](#)
6. [Distributed Systems](#)
7. [Exponential Backoff](#)
8. [Scalability](#)
9. [Tools For Scalability](#)
10. [Stateless API and Stateless Services](#)
11. [Observability](#)
12. [Reliability](#)
13. [Replication](#)
14. [Federation](#)
15. [FailOver](#)
16. [Graceful Degradation](#)
17. [SLOs/SLAs/SLIs](#)

18. Performance
19. Concurrency & Parallelism
20. Profiling and Optimization
21. Security
22. Partitioning / Sharding Data
23. Performance vs scalability
24. Consistent Hashing
25. Monolithic And Microservice Architecture
26. Optimistic vs Pessimistic Locking
27. Strong vs Eventual Consistency
28. Relational DB vs NoSQL
29. Types of NoSQL
30. Caching (Cache-aside, Write-through, Write-behind (Write-back), Refresh Ahead)
31. Cache Invalidation
32. Cache Eviction (LRU, LFU, MRU)
33. Data Center / Racks / Hosts
34. CPU / Memory / Hard Drive / Network Bandwidth
35. Random vs Sequential Read/Write on Disk
36. HTTP vs HTTP2 vs WebSockets
37. TCP/IP Model
38. IPv4 vs IPv6
39. TCP vs UDP
40. DNS Lookup
41. HTTPS & TLS
42. Public Key Infrastructure & Certificate Authority
43. Symmetric vs Asymmetric Key
44. Load Balancer → L4 vs L7
45. CDNs & Edge
46. Pull Vs Push CDNs
47. Bloom Filters & Count-min Sketch

48. Paxos – Consensus over Distributed Hosts
49. Leader Election
50. Virtual Machines & Containers
51. HyperVisor
52. Publisher–Subscriber or Queue
53. Fan-Out vs Fan-In
54. Map Reduce
55. Multithreading, Concurrency, Locks, Synchronization, CAS
56. Garbage Collection
57. Idempotency
58. Reverse Proxy
59. Network Protocols: TCP, WebSocket, HTTP, etc.
60. Client-Server vs Peer-to-Peer Architecture
61. Scale from 0 to Million Users
62. Decentralized Systems
63. Kafka
64. Message Queues
65. Proxy Servers
66. Storage Types
67. File System
68. Bloom Filter
69. Merkle Tree, Gossiping Protocol
70. Back of the envelope estimation (Storage Estimation, RPS Estimation, Bandwidth Estimation)
71. Sharding (Horizontal and Vertical)
72. Partitioning
73. Replication, Mirroring
74. Leader Election
75. Indexing etc.
76. NGINX
77. Load Balancer

- 73. Circuit Breaker
- 74. BulkHead Pattern
- 75. Retry
- 76. Polling
- 77. Load Balancing Algorithms
- 78. Least Privilege and Zero Trust Architecture
- 79. Fallback Patterns And Fallback Strategies:
- 30. What if redis goes down in the system?
- 31. gRPC Vs REST
- 32. Concurrency Patterns and Concurrency Scenarios
- 33. Federation
- 34. BackPressure
- 35. SQL tuning

- 1. Vertical vs Horizontal scaling

Vertical Scaling (Scaling Up)

Definition: Adding more power (CPU, RAM, SSD) to a single machine.

-  Example: Upgrading a server from 8 GB RAM to 64 GB RAM.
-  Pros:
 - Easier to implement (just upgrade hardware).
 - Requires less complex code and architecture.
-  Cons:
 - Has **hardware limits** – you can't scale infinitely.
 - If it fails, everything goes down (single point of failure).
 - Downtime might be needed to scale.

Horizontal Scaling (Scaling Out)

Definition: Adding more machines (nodes/servers) to your system.

-  Example: Moving from 1 server to a distributed cluster of 5 servers.
-  Pros:

- Virtually **infinite scaling** (add more machines as needed).
- Better **fault tolerance** (if one node fails, others continue).
- Supports distributed systems (microservices, load balancing, etc.)
- ! Cons:
 - More complex (load balancing, data partitioning, etc.)
 - Requires careful design for consistency and coordination.

| Scenario | Vertical | Horizontal |
|---------------------------------|----------|------------|
| Quick performance boost | ✓ | ✗ |
| Large-scale distributed system | ✗ | ✓ |
| Budget-limited startup | ✓ | ✗ |
| High availability & scalability | ✗ | ✓ |

Here's an easy-to-remember **pizza shop analogy**  to understand **Vertical vs Horizontal Scaling**:

Vertical Scaling = Making a Bigger Pizza Oven

Imagine you run a pizza shop. You have:

- **1 pizza oven**, and
- you're getting more orders.

So you:

- Replace your small oven with a **bigger, faster oven** that can bake more pizzas at once.
-  **Pros:** Simple! Just upgrade your oven.
-  **Cons:** One oven can only get so big. If it breaks, no pizzas at all.
-  This is **vertical scaling** – scaling up the power of your existing machine.

Horizontal Scaling = Adding More Pizza Ovens

Instead of upgrading your single oven:

- You **buy 5 more ovens**, hire a few chefs,
- and divide the orders among all ovens.

So:

- More pizzas baked **at the same time**.
- Even if one oven breaks, others keep running.

- **Pros:** Handles huge demand, highly fault-tolerant.
 - **Cons:** Requires coordination (who takes which order?), more space, and cost.
 - ➡ This is **horizontal scaling** – scaling *out* by adding more machines.
-

Summary in Analogy:

| Feature | Vertical Scaling | Horizontal Scaling |
|---|------------------------------|--|
|  Analogy | Bigger pizza oven | More pizza ovens |
|  Upgrade type | More power to 1 server | More servers |
|  Failure impact | Total outage if server fails | Minor impact; others can still work |
|  Complexity | Low | High (needs coordination/load balancing) |
|  Scalability limit | Yes (hardware limit) | No (theoretically infinite) |

2. **CAP theorem**
 - CAP** theorem is a principle that applies to **distributed systems**, stating that a system can satisfy **only two out of the three** guarantees at any given time:
 - Consistency (C)** – Every read receives the most recent write or an error.
 - Availability (A)** – Every request receives a (non-error) response, without guarantee that it contains the most recent write.
 - Partition Tolerance (P)** – The system continues to operate despite network partitions (failures in communication between nodes).

-  You **must** tolerate partitions in distributed systems, so usually it's a trade-off between:
- **CP** (e.g., HBase, MongoDB when configured for consistency)
 - **AP** (e.g., Cassandra, DynamoDB)

ACID properties apply to **database transactions** (particularly in relational databases):

1. **Atomicity** – All parts of a transaction succeed or none at all.
2. **Consistency** – The database moves from one valid state to another, maintaining rules like constraints, triggers, cascades, etc.
3. **Isolation** – Transactions do not interfere with each other.
4. **Durability** – Once committed, data changes are permanent, even in case of a crash.

 Here, **Consistency** refers to **database integrity**, ensuring that the **schema rules**, **foreign keys**, and **constraints** are not violated.

In Summary:

| | | |
|------------|----------------------------------|--|
| Aspect | CAP Consistency | ACID Consistency |
| Applies to | Distributed systems | Database transactions |
| Meaning | All nodes see the same data | All data remains valid under all rules |
| Focus | Synchronization between replicas | Integrity of a single database |

3. ACID vs BASE

◆ ACID (Traditional Databases like RDBMS)

ACID guarantees strong consistency and reliability.

| Property | Meaning |
|------------------------|--|
| A - Atomicity | All operations in a transaction either complete or none do. |
| C - Consistency | The database moves from one valid state to another valid state. |
| I - Isolation | Concurrent transactions do not interfere with each other. |
| D - Durability | Once a transaction is committed, it remains so, even in case of crash. |

➡ Used in: SQL databases (MySQL, PostgreSQL, Oracle)

◆ BASE (Modern NoSQL/Distributed Systems)

BASE trades strict consistency for availability and performance.

| Property | Meaning |
|---------------------------------|---|
| B - Basically Available | System guarantees availability (some data may be outdated). |
| S - Soft state | System state may change over time, even without input. |
| E - Eventual consistency | System becomes consistent over time, given no new updates. |

➡ Used in: NoSQL databases (Cassandra, DynamoDB, Couchbase)

4. Fault Tolerance

Fault tolerance is a system's ability to continue operating properly even if some part of it fails.

◆ Why it Matters

In real-world distributed systems, **failures are inevitable** (network issues, hardware crashes, power outages). A fault-tolerant system **detects** these failures and **recovers gracefully** without affecting the user experience.

◆ Types of Faults

| Type | Example |
|-------------------------|---------------------------|
| Hardware Failure | Disk crash, power failure |

| | |
|------------------------|-----------------------------------|
| Network Failure | Packet loss, server unreachable |
| Software Bugs | Exceptions, memory leaks |
| Human Errors | Wrong config, accidental deletion |

◆ Techniques to Achieve Fault Tolerance

| Technique | Description |
|--------------------------------|---|
| Replication | Duplicate data/services across multiple machines |
| Failover | Automatically switch to a backup component on failure |
| Redundancy | Have extra hardware/software in place as a backup |
| Retry Logic | Retry failed operations with exponential backoff |
| Circuit Breaker Pattern | Stop sending requests to a failing service temporarily |
| Load Balancing | Distribute requests across multiple servers |
| Data Backups | Regular snapshots of data to restore in case of failure |

◆ Real-World Example

If an Amazon data center goes down, you can still shop online because AWS automatically reroutes traffic to healthy zones. That's fault tolerance in action.

Fault-tolerant architectures like **Raft**, **Paxos**, or real use cases from **Netflix**, **Google**, etc.

5. **Distributed Systems:**

🌐 What is a Distributed System?

A **distributed system** is a collection of independent computers that appears to its users as a **single coherent system**.

🧠 Key Characteristics

| Feature | Description |
|-------------------------|--|
| Concurrency | Many components can operate simultaneously. |
| Scalability | Can handle growth in users, data, or traffic by adding more nodes. |
| Fault Tolerance | Continues to work even if parts of the system fail. |
| Transparency | Hides the complexity from users (location, replication, failures). |
| Decentralization | No single point of control or failure. |

🔧 Core Components

- **Nodes** – Individual machines/servers
 - **Message Passing** – Communication between nodes (typically via HTTP, gRPC, etc.)
 - **Consistency Mechanisms** – Like Quorum, Raft, or Paxos
 - **Coordination Services** – e.g., **ZooKeeper**, **etcd**
-

Common Challenges

| Challenge | Description |
|------------------------------|--|
| Network Latency | Messages between nodes are slower than in-memory. |
| Partial Failures | Some nodes may go down while others work fine. |
| Data Consistency | Keeping data in sync across all nodes is tricky. |
| Clock Synchronization | No global clock; causes problems in ordering events. |
| Split-Brain Scenarios | Multiple leaders may be elected accidentally. |

Common Concepts

- **CAP Theorem**: Consistency, Availability, Partition Tolerance – Pick any 2.
 - **ACID vs BASE**: Trade-offs between strict and eventual consistency.
 - **Leader Election**: Choosing a master node (e.g., via Raft).
 - **Gossip Protocol** : Used for peer discovery and state sharing.
 - **Replication & Sharding**: For scaling and fault tolerance.
-

Real-World Examples

- **Google File System (GFS)**
- **Amazon DynamoDB**
- **Apache Cassandra**
- **Kubernetes (Distributed Orchestration)**
- **Bitcoin/Ethereum (Blockchain)**

Let me know if you'd like a **visual mind map**, or want me to relate this to **interview system design questions** like:

- "Design a URL Shortener"

- "Design Dropbox"
 - "Design WhatsApp"
-

6. **Exponential Backoff**

What is Exponential Backoff?

Exponential Backoff is a retry strategy where a system **waits for increasingly longer intervals** before retrying a failed operation. It's commonly used in:

- API requests
 - Network communication
 - Distributed databases
-

How it works:

After a failure, wait time grows **exponentially** with each retry.

Example (base delay = 1s):

```
Retry 1 → wait 1s  
Retry 2 → wait 2s  
Retry 3 → wait 4s  
Retry 4 → wait 8s  
...and so on
```

Formula:

```
wait_time = base * (2^retry_count)
```

Optional Enhancements:

| Enhancement | Purpose |
|----------------------|---|
| Jitter | Adds randomness to avoid retry storms (e.g., 4s ± random ms). |
| Max Wait Time | Puts a cap on the maximum delay. |
| Max Retries | Limits total number of retries. |

Where It's Used:

- AWS SDKs

- Google Cloud APIs
 - Databases like DynamoDB, Bigtable
 - HTTP `429 (Too Many Requests)` handling
 - Retrying Kafka producers or consumers
-

Benefits

- Reduces **network congestion**
- Helps systems **recover gracefully**
- Avoids `retry storm`s in distributed environments

7. `Scalability`

Scalability refers to a system's ability to handle **increased load** (more data, users, or requests) by **adding resources** (hardware or software) — without sacrificing performance.

Two Main Types of Scalability:

◆ 1. Vertical Scalability (Scaling Up)

Add more resources **to a single machine** (CPU, RAM, SSD).

-  Easier to implement
-  Limited by hardware ceiling
-  Example: Upgrade your database server from 8GB RAM to 64GB

◆ 2. Horizontal Scalability (Scaling Out)

Add **more machines/nodes** to distribute the load.

-  Can scale almost infinitely
 -  Harder to manage (requires distributed systems)
 -  Example: Add more web servers behind a load balancer
-

Scalability in Different Layers

| Layer | Scalability Strategy |
|----------|---|
| Frontend | CDN, caching, static content |
| Backend | Microservices, load balancing, stateless APIs |
| Database | Read replicas, sharding, NoSQL |

| | |
|---------------------|--|
| Cache | Redis/Memcached clustering |
| Queue | Kafka, RabbitMQ for async processing |
| File Storage | S3 or distributed file systems (Ceph, GFS) |

Key Techniques

- **Load Balancing**
- **Database Sharding**
- **Replication**
- **Caching (Redis, CDN)**
- **Stateless Services**
- **Asynchronous Processing (Queues)**
- **Auto-scaling (e.g., AWS Auto Scaling Groups)**

How to Measure Scalability?

- **Throughput:** Requests/sec
- **Latency:** Response time
- **Resource Utilization:** CPU, memory
- **Cost efficiency:** How well it scales with cost

8. [Tools For Scalability](#) (Container Orchestration And Database Scaling)

Infrastructure & Orchestration

| Tool | Purpose |
|--|---|
| Kubernetes | Orchestrates containerized applications, auto-scales pods based on demand |
| Docker | Containerization for lightweight, consistent deployments |
| Terraform / Pulumi | Infrastructure as Code (IaC) for scalable cloud resource provisioning |
| Amazon EC2 Auto Scaling / GCP Managed Instance Groups | Automatically adjusts compute resources based on load |
| Nomad | Lightweight, simpler orchestration alternative to Kubernetes |

Load Balancing & Traffic Management

| Tool | Purpose |
|------------------------------------|--|
| NGINX / HAProxy | Reverse proxy and load balancing for web traffic |
| Envoy | High-performance edge/microservice proxy for service mesh |
| AWS ALB / GCP Load Balancer | Managed cloud load balancing at scale |
| Istio / Linkerd | Service mesh for advanced traffic routing, security, and observability |

Caching & Fast Data Access

| Tool | Purpose |
|----------------------------------|---|
| Redis / Memcached | In-memory key-value store for reducing DB load |
| CDNs (Cloudflare, Akamai) | Cache static content close to users, reduce latency |
| Varnish Cache | HTTP accelerator for content-heavy websites |

Asynchronous Processing & Messaging

| Tool | Purpose |
|------------------------------------|---|
| Kafka | Distributed message queue for high-throughput event streaming |
| RabbitMQ / NATS | Message brokers for decoupling components |
| Amazon SQS / Google Pub/Sub | Scalable, managed message queues |

Scalable Databases

| Tool | Purpose |
|--------------------------------|---|
| Cassandra / ScyllaDB | Distributed NoSQL DBs designed for horizontal scaling |
| MongoDB Atlas | Scalable document database with sharding |
| Amazon Aurora / Spanner | Managed relational DBs with high availability and scaling |
| Elasticsearch | Scalable search and analytics engine |

Observability & Monitoring

| Tool | Purpose |
|-----------------------------|--------------------------------------|
| Prometheus + Grafana | Metrics collection and visualization |

| | |
|---------------------------------|---|
| Datadog / New Relic / Dynatrace | Full-stack observability for infra + apps |
| Jaeger / Zipkin | Distributed tracing for microservices |

🧪 Testing for Scalability

| Tool | Purpose |
|---------------|---|
| Apache JMeter | Load testing for web apps/services |
| k6 | Modern load testing tool with scripting |
| Locust | Python-based load testing framework |

9. Stateless API and Stateless Services

📌 Stateless APIs

- An **API is stateless** when **each HTTP request** contains all the information the server needs to understand and process it.
- The server **doesn't retain** any session, history, or user state between requests.

✅ Example:

```
GET /user/123
Authorization: Bearer <token>
```

- The token contains all necessary info (e.g., user identity, roles).
- Server reads the token, processes the request, and **forgets** everything after sending the response.

🔒 Benefits:

- Easy to scale horizontally (no session stickiness)
- Cache-friendly
- Resilient (no dependency on shared memory/session)

📌 Stateless Services

- A **stateless service** is a broader concept where the **service itself doesn't maintain state** between any interactions.
- Applies to **all types of service logic**, not just API layers (e.g., background workers, microservices).

Example:

- A payment processing service that receives an order ID, fetches order info from a database, and completes the transaction — all without maintaining any session context.
-

Comparison

| Feature | Stateless API | Stateless Service |
|-------------------|----------------------------------|---|
| Scope | Communication layer (e.g., REST) | Whole service (incl. logic, processing) |
| Focus | HTTP request/response | Service logic and processing |
| State persistence | None per request | None between invocations |
| Example | RESTful endpoint | Microservice for processing orders |
| Benefit | Easy to scale, secure | Easy to restart, distribute |

Where State Goes Instead?

- Client (e.g., browser or mobile app)
- Databases
- Caches (e.g., Redis)
- Tokens (e.g., JWT)
- External storage

9. Observability

Observability is the ability to understand what's happening inside a system **based on external outputs** — such as logs, metrics, and traces. It helps developers and operators **detect, diagnose, and resolve issues** quickly in complex, distributed systems.

Core Pillars of Observability

| Pillar | Description |
|----------------|--|
| Logs | Time-stamped records of events (errors, info, debug messages). |
| Metrics | Numeric data points representing system health (CPU usage, error rates, etc.). |
| Traces | End-to-end journey of a request across services/components. |

Example

If an e-commerce site slows down:

- **Logs** help you find errors (e.g., "DB timeout").

- **Metrics** show a spike in DB latency.
 - **Traces** reveal which microservice caused the delay.
-

Why Observability Matters

- Helps in **root cause analysis**
 - Essential for **debugging** distributed systems
 - Enables **proactive monitoring**
 - Enhances **incident response** and **MTTR** (Mean Time To Recovery)
 - Aids in **performance optimization**
-

Common Observability Tools

| Tool | Purpose |
|--|--------------------------|
| Prometheus | Metrics collection |
| Grafana | Metrics visualization |
| ELK Stack (Elasticsearch, Logstash, Kibana) | Logging |
| Fluentd / Logstash | Log shipping |
| Jaeger / Zipkin | Distributed tracing |
| Datadog / New Relic / Splunk | All-in-one observability |

Best Practices

- Correlate logs, metrics, and traces with a common ID (e.g., request ID)
- Set up alerts on critical metrics (CPU, memory, error rates)
- Use **structured logging** (JSON) for easy parsing
- Monitor **SLOs/SLAs**
- Always track **latency, traffic, errors, and saturation (L-T-E-S)**

10. **Reliability**

Reliability in a system refers to **its ability to consistently perform its intended function without failure** over time. In distributed systems, a **reliable system continues to work correctly, even in the face of hardware failures, network issues, or software bugs**.

Key Concepts of Reliability

1. Redundancy

- **Replicate critical components** (servers, databases, load balancers).
- Example: Master-slave DB replication; multiple availability zones.

2. Fault Tolerance

- System can handle **hardware/software failures** without crashing.
- Often achieved via techniques like **failover, circuit breakers, and retry mechanisms**.

3. Failover & Backup

- Automatically switch to a **redundant system** when a failure is detected.
- Use **database snapshots, incremental backups, and geo-replication**.

4. Monitoring and Alerting

- Tools like **Prometheus, Grafana, Datadog, or ELK Stack**.
- Helps detect and respond to failures **proactively**.

5. Load Balancing

- Prevents any one node from getting overwhelmed.
- Example: **NGINX, AWS ELB, Cloudflare Load Balancing**.

6. Graceful Degradation

- System continues to offer **partial functionality** instead of failing completely.
- Example: A video site can show previews if streaming fails.

7. Retries with Exponential Backoff

- Handle transient failures by retrying requests with delay.
- Prevents overwhelming failing services.

Measuring Reliability

| Metric | Description |
|-------------------------------------|--|
| Uptime | % of time the system is operational (e.g., 99.99%) |
| MTTF (Mean Time To Failure) | Average time between failures |
| MTTR (Mean Time To Recovery) | Time taken to fix a failure |

| | |
|-------------|---|
| SLOs | Targets like "Service should respond in <300ms 99.9% of the time" |
| SLAs | Formal agreements (often with clients) based on SLOs |

✓ Tools & Practices for High Reliability

| Area | Tools / Practices |
|------------------------|--|
| Backups | AWS RDS Snapshots, rsync, Velero (for Kubernetes) |
| Monitoring | Prometheus, Grafana, Datadog, CloudWatch |
| Resilience | Hystrix, Resilience4j, Circuit Breakers |
| Testing | Chaos Engineering (e.g., Netflix Chaos Monkey) |
| CI/CD | Automated pipelines reduce human error and allow safer rollbacks |
| Error Budgeting | Balance between reliability and innovation speed |

11. SLOs/SLAs

📊 SLOs vs SLAs vs SLIs (The Golden Trio of Reliability)

These terms are often used in the context of **reliability engineering** and **observability** — especially in production-grade, cloud-native systems.

🔄 Quick Definitions

| Term | Full Form | What it Means |
|------------|-------------------------|---|
| SLA | Service Level Agreement | A formal contract between a provider and customer defining guaranteed levels of service. |
| SLO | Service Level Objective | An internal goal set by the service team to maintain performance and reliability. |
| SLI | Service Level Indicator | A measurable metric (like uptime, latency) used to track how the service is performing. |

🧩 Relationship Between the Three

- **SLIs** are the metrics you track.
- **SLOs** are the reliability targets you aim to meet based on SLIs.
- **SLAs** are what you promise to your users/customers (and may include penalties if violated).

✓ Example

| Concept | Value |
|------------|---|
| SLI | 99.95% of HTTP requests respond within 300ms |
| SLO | Maintain 99.9% latency under 300ms over a 30-day window |
| SLA | We guarantee 99.5% uptime monthly; failure results in 10% credit refund |

💥 If Things Go Wrong?

- **Breach of SLO:** Internal signal to improve reliability or reduce new feature rollout.
- **Breach of SLA:** May trigger customer compensation, financial penalties, or reputational damage.

📌 Why SLOs Matter More for Engineers

SLOs help:

- Balance reliability vs velocity (error budgets).
- Make **data-driven decisions** about deployments.
- Avoid overengineering — you don't need 100% reliability if your users are fine with 99.9%.

📊 Common SLIs You Might Track

| SLI | Description |
|---------------------|---|
| Availability | % of successful requests |
| Latency | % of requests below a certain response time |
| Error Rate | % of failed requests |
| Throughput | Requests per second handled |
| Durability | Data not lost over time |

12. Performance

Performance refers to **how efficiently a system handles workload**, including how fast it responds to requests, how much it can process concurrently, and how it scales under load.

🚀 Key Performance Metrics

| Metric | Meaning |
|----------------|---|
| Latency | Time taken to process a request (usually in ms) |

| | |
|---------------------------------|--|
| Throughput | Number of requests processed per unit time (req/sec) |
| QPS (Queries Per Second) | Specific throughput for read-heavy APIs or databases |
| Concurrency | Number of simultaneous operations handled |
| Response Time | End-to-end time a user waits for a response |
| Resource Utilization | CPU, memory, disk, and network usage efficiency |

🛠️ Tools for Performance Monitoring

| Layer | Tools |
|-----------------|--|
| Frontend | Lighthouse, Chrome DevTools |
| Backend | JMeter, Apache Bench, wrk, Locust |
| Infra | Grafana + Prometheus, Datadog, New Relic |
| Tracing | Jaeger, Zipkin (via OpenTelemetry) |

🧠 Common Bottlenecks

- Slow **database queries** (missing indexes, joins)
- **Synchronous dependencies** (waiting on 3rd party APIs)
- **Poor caching** strategy
- **Thread starvation** or high **GC overhead**
- Network **latency or packet loss**

✅ Best Practices to Improve Performance

1. Caching

- Use **Redis, Memcached, browser caches, CDNs** (Cloudflare, Akamai).

2. Load Balancing

- Distribute traffic using **Round Robin, Least Connections, or IP Hashing**.

3. Async Processing

- Use **message queues** like Kafka, RabbitMQ, SQS for long-running tasks.

4. Efficient Data Access

- Use **database indexing**, avoid N+1 queries, use connection pools.

5. Compression

- GZIP responses, minify JS/CSS, use Brotli for better compression.

6. Horizontal Scaling

- Add more machines/containers instead of overloading one.

7. Concurrency & Parallelism

- Use thread pools , async/await , coroutines to handle tasks concurrently.

8. Profiling and Optimization

- Use tools like YourKit , VisualVM , perf , flamegraphs to identify hot spots.

🧪 Performance Testing Types

| Type | Purpose |
|----------------|--|
| Load Testing | Checks system under expected load |
| Stress Testing | Push beyond limits to see how it fails |
| Spike Testing | Sudden burst of traffic |
| Soak Testing | Long-duration load to test memory leaks, degradation |

13. Concurrency & Parallelism

| | |
|-------------|--|
| Concurrency | The ability to handle multiple tasks at once, by switching between them. It gives an illusion of simultaneity. |
| Parallelism | The ability to execute multiple tasks at the same time , using multiple processors or cores . |

⚙️ Analogy

Imagine you're cooking:

- **Concurrency:** You're cooking pasta, and while it's boiling, you chop veggies. You're managing multiple tasks by **switching** between them — one at a time.
- **Parallelism:** You have a friend — while you cook pasta, they chop veggies at the **same time**.

🔧 Tech Stack Context

| Feature | Concurrency | Parallelism |
|---------|-------------|-------------|
| | | |

| | | |
|------------------|--|---|
| CPU Cores Needed | 1+ | 2+ |
| Example | Async I/O, threads | Multi-threaded matrix operations |
| Java Example | <code>ExecutorService</code> , <code>CompletableFuture</code> | <code>ForkJoinPool</code> , <code>parallelStream()</code> |
| JavaScript | Event loop, Promises (<code>async/await</code>) | Web Workers (browser), Worker Threads (Node.js) |
| Python | <code>asyncio</code> , <code>threading</code> | <code>multiprocessing</code> |
| Use case | I/O-bound tasks | CPU-bound tasks |

💡 When to Use What?

| Type of Task | Use | Reason |
|----------------------|--------------------|------------------------|
| API Calls, File I/O | Concurrency | Latency due to waiting |
| Image processing, ML | Parallelism | Heavy CPU usage |

🚦 Example in Java

```
ExecutorService executor = Executors.newFixedThreadPool(4);

executor.submit(() -> {
    System.out.println("Running in parallel!");
});
```

For concurrency using `CompletableFuture` :

```
CompletableFuture.runAsync(() -> {
    System.out.println("Concurrent task running");
});
```

🧩 TL;DR

- **Concurrency:** Managing many tasks at once.
- **Parallelism:** Running many tasks at the same time.
- **You can have concurrency without parallelism**, and vice versa.

14. Profiling and Optimization

Profiling and **Optimization** are crucial stages in improving the **performance** and **efficiency** of an application.

What is Profiling?

Profiling is the process of analyzing a program to understand how resources (CPU, memory, threads, etc.) are being used.

Common Profiling Metrics:

- **CPU Usage:** Which functions/methods consume the most CPU time?
- **Memory Allocation:** Where is memory allocated and possibly leaked?
- **Thread Activity:** Are there deadlocks, race conditions, or thread contention?
- **I/O Wait Times:** How much time is spent waiting on I/O?

Tools for Profiling

| Language/Platform | Tools |
|-------------------|--|
| Java | VisualVM, YourKit, JProfiler, Java Flight Recorder (JFR) |
| JavaScript | Chrome DevTools, Lighthouse |
| Python | cProfile, memory_profiler |
| Node.js | Clinic.js, Node Profiler |
| System-wide | <code>perf</code> , <code>htop</code> , <code>dstat</code> , <code>valgrind</code> , <code>strace</code> |

What is Optimization?

Optimization is the act of modifying a system to make it more efficient — faster, smaller, more responsive, or less resource-hungry.

Optimization Loop:

1. **Profile** to identify bottlenecks.
2. **Analyze** what's consuming the most resources.
3. **Optimize** by refactoring code, changing algorithms, reducing memory usage, etc.
4. **Benchmark** to verify improvement.
5. Repeat if needed.

Types of Optimization

| Type | Example |
|-------------------|---|
| Code-level | Avoiding nested loops, memoization, caching |

| | |
|--------------------|--|
| Algorithmic | Replacing $O(n^2)$ with $O(n \log n)$ |
| Memory | Reusing buffers, reducing object creation |
| Concurrency | Using thread pools, async patterns |
| Database | Indexing, avoiding N+1 queries, query optimization |
| Network | Compression, batching requests, CDN |

Example in Java (JFR)

```
# Start your app with:  
java -XX:StartFlightRecording:filename=recording.jfr MyApp
```

Then open `.jfr` in **Java Mission Control** to analyze performance.

Tips for Effective Profiling & Optimization

- Profile **before** optimizing — don't guess
- Optimize **hot paths**, not everything
- Focus on **95th percentile** latency, not just average
- Use **load testing tools** like JMeter or Locust
- Know your **resource constraints**

Java Profiling Tools

| Tool | Features |
|-----------------------------------|--|
| VisualVM | Bundled with JDK, real-time monitoring, memory analysis |
| Java Flight Recorder (JFR) | Low-overhead profiling, production-ready, integrated with JVM |
| JProfiler | Commercial, deep profiling (CPU, memory, threads, JDBC, JPA, etc.) |
| YourKit | Commercial, user-friendly, supports integration with IDEs |
| Async Profiler | Low-overhead sampling profiler for CPU and memory |

JavaScript / Frontend

| Tool | Features |
|------------------------|---|
| Chrome DevTools | Performance tab: frame rate, paint time, JS execution |
| Lighthouse | Web performance scores, audits accessibility, SEO |
| WebPageTest | Analyzes performance from multiple locations/devices |

Node.js

| Tool | Features |
|--------------------------|--|
| Node.js --inspect | Built-in inspector to connect with Chrome DevTools |
| Clinic.js | Suite (Doctor, Flame, Bubbleprof) for CPU/memory profiling |
| v8-profiler | CPU and heap snapshots, flamegraphs |

Python

| Tool | Features |
|------------------------|---|
| cProfile | Built-in deterministic profiler for CPU usage |
| memory_profiler | Line-by-line memory usage |
| Py-Spy | Sampling profiler, no code changes, works in production |
| line_profiler | Time per line of code, great for algorithm bottlenecks |

System-level Profilers

| Tool | Platform | Features |
|------------------------|-------------|--|
| perf | Linux | Event-based sampling profiler |
| htop / atop | Linux | CPU/memory usage visualization |
| strace / dtrace | Linux/macOS | Trace system calls and signals |
| valgrind | Linux | Memory leaks, heap usage, CPU cycles |
| gprof | GNU | Call graph and flat profile for C/C++ programs |

Database Profiling

| DB | Tool |
|-------------------|--|
| MySQL | <code>EXPLAIN</code> , <code>SHOW PROFILE</code> , Percona Toolkit |
| PostgreSQL | <code>EXPLAIN ANALYZE</code> , <code>pg_stat_statements</code> |
| MongoDB | <code>db.setProfilingLevel</code> , Atlas Performance Advisor |

Cloud/Infra Profilers

| Platform | Tool |
|--------------------------------------|--|
| AWS | AWS X-Ray, CloudWatch Profiler |
| GCP | Cloud Profiler (low overhead, production safe) |
| Datadog, New Relic, Dynatrace | End-to-end distributed tracing and profiling |

1. Code-Level Optimization Tools

These help identify bottlenecks in code and suggest improvements.

| Tool | Language | Purpose |
|--|-----------------------|--|
| JMH (Java Microbenchmark Harness) | Java | Benchmark method-level performance |
| GraalVM | Java | High-performance JIT compiler |
| ESLint / TSLint | JavaScript/TypeScript | Detect inefficient or bad code practices |
| Blackfire | PHP | Profiling and performance optimization |
| Py-Spy / line_profiler | Python | CPU/memory/time usage per line |
| Valgrind | C/C++ | Memory leaks, cache use, branch prediction |
| Go tool pprof | Go | CPU, memory, and goroutine profiling |

🧠 2. Compiler-Level Optimization Tools

These tools optimize code at compile-time.

| Tool | Language | Optimization Type |
|--------------------------|----------------|---|
| GCC/Clang -O2/-O3 | C/C++ | Compiler flags for loop unrolling, inlining |
| R8 / ProGuard | Java (Android) | Code shrinking, obfuscation |
| Babel / SWC | JavaScript | Transpilation and dead-code elimination |

⚙️ 3. Runtime Optimization Tools

Tools that help improve runtime performance.

| Tool | Platform | Features |
|------------------------|----------|--|
| JVM GC tuning | Java | Reduce latency, optimize memory |
| Node.js Cluster | Node.js | Multi-core load balancing |
| NGINX | Web | Reverse proxy, caching, gzip compression |
| Varnish | Web | HTTP caching to reduce backend load |

📊 4. Application Performance Monitoring (APM)

Used to monitor, trace, and optimize applications in production.

| Tool | Features |
|------------------|--|
| Datadog | Real-time metrics, traces, dashboards |
| New Relic | Code-level traces, alerts, database optimization |

| | |
|--------------------|---|
| Dynatrace | AI-driven performance analytics |
| AppDynamics | Transaction tracing, resource bottlenecks |

5. Cloud & Infra Optimization Tools

Optimizing usage of cloud resources and infra.

| Tool | Cloud | Features |
|----------------------------------|-------|--|
| AWS Trusted Advisor | AWS | Cost, performance, security optimization |
| GCP Recommender | GCP | Suggests scaling, resource changes |
| Kubernetes Metrics Server | K8s | Autoscaling decisions |
| Cloudflare | CDN | Performance & edge optimizations |

6. Frontend Optimization Tools

Improve performance of web apps.

| Tool | Purpose |
|--------------------------------------|--|
| Lighthouse | Audits performance, accessibility, SEO |
| Webpack Bundle Analyzer | Analyze and reduce JS bundle size |
| ImageOptim / Squoosh | Image compression |
| Lazy loading + Code splitting | Native techniques in React/Angular/Vue |

7. Database Optimization

Tools to optimize query execution and DB performance.

| Tool | DB | Purpose |
|--------------------------|------------------|----------------------------------|
| EXPLAIN / ANALYZE | MySQL/PostgreSQL | Query plan analysis |
| pt-query-digest | MySQL | Identify slow queries |
| Indexes / Caching | Any | Improve data access |
| Redis / Memcached | Any | Offload frequently accessed data |

15. Security

1. Core Security Principles (CIA Triad)

| Principle | Description |
|------------------------|---|
| Confidentiality | Ensures data is accessible only to authorized users. (e.g., encryption, access control) |
| Integrity | Ensures data is not tampered with. (e.g., hashing, checksums) |

| | |
|---------------------|--|
| Availability | Ensures systems are available when needed. (e.g., redundancy, DDoS protection) |
|---------------------|--|

2. Key Areas in Software Security

Authentication & Authorization

- *Authentication*: Verify user identity (e.g., OAuth, JWT, biometric).
- *Authorization*: Define what users can access (e.g., RBAC, ABAC).

Data Protection

- Use HTTPS (TLS)
- At-rest encryption (e.g., AES-256)
- In-transit encryption (TLS 1.2+)

Input Validation & Sanitization

- Prevent SQL Injection, XSS, CSRF
- Use whitelists, parameterized queries

Secure APIs

- Rate limiting
- OAuth2 / API key usage
- Prevent sensitive data leakage

Secure Storage

- Never store plaintext passwords (use bcrypt, Argon2)
- Secure secrets management (e.g., HashiCorp Vault, AWS Secrets Manager)

3. Testing and Monitoring Tools

| Tool | Purpose |
|----------------------------------|--|
| OWASP ZAP / Burp Suite | Penetration testing |
| SonarQube / Semgrep | Static code analysis |
| Snyk / Dependabot | Vulnerability scanning in dependencies |
| Auditd / Falco | System-level intrusion detection |
| WAF (Cloudflare, AWS WAF) | Application-layer firewall protection |

4. Security by Design

- **Zero Trust Architecture** – “Never trust, always verify”
 - **Least Privilege Principle** – Only give minimum required access
 - **Defense in Depth** – Multiple layers of security (e.g., firewall + IDS + encryption)
-

5. Standards and Frameworks

- **OWASP Top 10** – Awareness of top vulnerabilities
 - **NIST Cybersecurity Framework**
 - **ISO/IEC 27001** – InfoSec Management
 - **SOC 2 / GDPR / HIPAA** – Compliance standards
-

6. DevSecOps

- Integrate security in the CI/CD pipeline:
 - Secret scanning before deploys
 - Linting for insecure patterns
- SAST (Static Analysis) + DAST (Dynamic Analysis)

1. Authentication & Authorization

Authentication (Who are you?)

- Use secure methods like **OAuth2**, **OIDC**, **SAML**, or **JWT**.
- Integrate with Identity Providers (e.g., **Auth0**, **Okta**, **Keycloak**).

Authorization (What can you do?)

- Use **RBAC** (Role-Based Access Control) or **ABAC** (Attribute-Based Access Control).
 - Centralize access policies using **Policy Engines** (e.g., **OPA - Open Policy Agent**).
-

2. Data Security

At Rest

- Encrypt sensitive data using **AES-256** or equivalent.
- Store secrets in a **secure vault** (e.g., **AWS Secrets Manager**, **HashiCorp Vault**).

In Transit

- Enforce **TLS/SSL** for all client-server & inter-service communication.
 - Use **mTLS** for internal service-to-service communication.
-

3. Secure APIs

- Authenticate all APIs.
 - Use **rate limiting** and **throttling** to prevent abuse.
 - Validate input (against **SQLi, XSS, CSRF**, etc.).
 - Don't expose internal implementation details.
-

4. Service Layer Security

- Enforce **Zero Trust Architecture**: "Never trust, always verify".
 - Services should **only** talk to the services they need (network segmentation).
 - Use **API gateways** to manage ingress and apply policies.
 - Use **service mesh** (e.g., Istio/Linkerd) to enforce security at runtime.
-

5. Logging, Auditing & Monitoring (Observability)

- **Log** authentication attempts, permission changes, data access.
 - Enable **audit trails** for sensitive operations.
 - Use tools like **ELK, Datadog, or Prometheus + Grafana** for alerts and dashboards.
 - Alert on anomalies and potential breaches.
-

6. Network Security

- Restrict traffic with **security groups, firewalls, and VPCs**.
 - Use **WAFs** (Web Application Firewalls) to block attacks at the edge.
 - DDoS Protection via services like **Cloudflare, AWS Shield, or Google Cloud Armor**.
-

7. CI/CD Security (DevSecOps)

- **Shift-left**: Scan dependencies early (e.g., Snyk, Dependabot).
- Lint and statically analyze code (e.g., SonarQube, Semgrep).
- Use **container image scanners** (e.g., Trivy, Clair).
- Secrets scanning (e.g., Gitleaks, GitGuardian).

- Enable **build-time & pre-deploy** security checks.
-

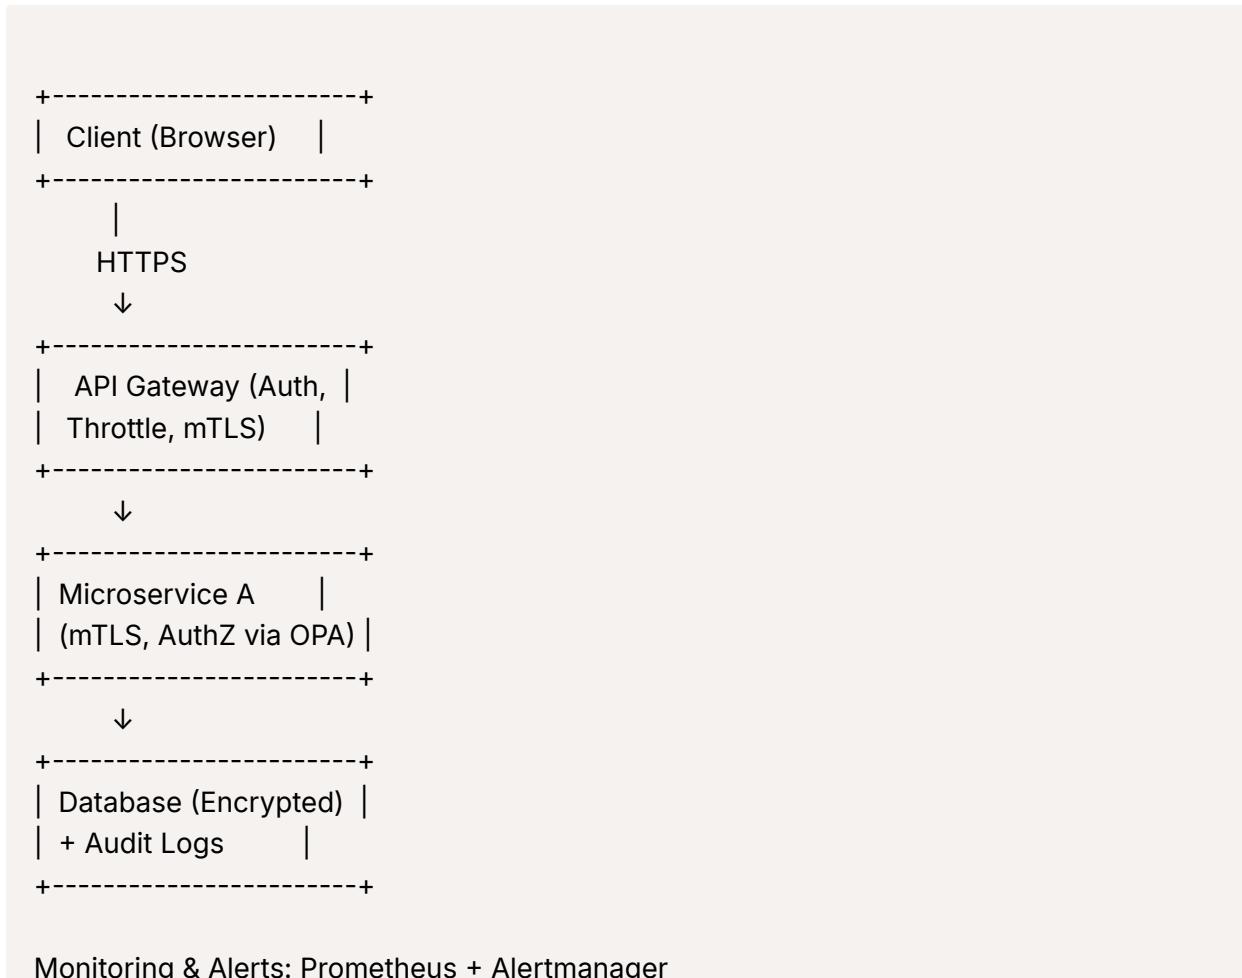
8. User-Facing App Security

- Secure cookies (`HttpOnly` , `Secure` , `SameSite`).
 - Enable **CSP** (Content Security Policy).
 - Prevent **Clickjacking** using `X-Frame-Options` .
 - Sanitize all inputs to prevent **XSS**, **CSRF**, etc.
-

9. Compliance & Governance

- Follow security frameworks: **OWASP**, **NIST**, **SOC 2**, **ISO 27001**, **PCI DSS**, etc.
 - Implement **data retention** and **data deletion** policies.
 - Apply **GDPR** or **HIPAA** if applicable.
-

Diagram Sample (Security Applied in System Design)



Secrets: HashiCorp Vault

Audit Trails: Fluentd + ElasticSearch

✓ Final Checklist for Secure Design

- Authentication & Authorization in place
- Encrypted communication (TLS/mTLS)
- Input validation and rate limiting
- Secure APIs & secrets management
- Continuous security monitoring
- Proper logging, audit, and alerting
- Least privilege + Zero Trust architecture

16. [Partitioning / Sharding Data](#)

◆ Partitioning vs Sharding

- **Partitioning** is the general practice of splitting a large dataset into smaller, more manageable pieces (partitions) based on some key.
- **Sharding** is a form of horizontal partitioning where each partition (shard) lives on its own database instance/server.

🧩 Why Partition / Shard?

1. **Performance:** Spreads reads/writes across multiple servers.
2. **Scalability:** You can add more shards as data grows.
3. **Availability:** Failure of one shard doesn't bring down the entire system.
4. **Manageability:** Smaller chunks are easier to backup, restore, and maintain.

🔧 Common Sharding Strategies

| Strategy | How It Works | Pros | Cons |
|------------------|---|---|---|
| Range | Assign contiguous key-ranges to shards. | Simple to understand; good for range scans. | Hotspots if data skews to one range. |
| Hash | Apply a hash function to the shard key. | Evenly distributes writes; avoids hotspots. | Harder to do range queries. |
| Directory | Maintain a lookup table: key → shard. | Fully flexible; can rebalance dynamically. | Directory is a single point of failure. |

| | | | |
|------------------------|--|---|--|
| Consistent Hash | Keys & shards mapped onto a hash ring. | Automatic rebalancing when shards join/leave. | More complex; virtual node management. |
|------------------------|--|---|--|

Design Considerations

1. Choosing Your Shard Key

- Should be:
 - **Highly cardinal** (many distinct values)
 - **Uniformly distributed** writes and reads
- Avoid: sequential or timestamp keys that produce hot shards.

2. Rebalancing & Resharding

- Plan how you'll move data when adding/removing shards.
- Use tools or orchestrated live migrations (e.g., MongoDB's balancer, Vitess resharding).

3. Cross-Shard Queries & Joins

- Minimize cross-shard joins—keep related data on the same shard.
- If unavoidable, aggregate at the application layer or use a middleware.

4. Transactions & Consistency

- Single-shard transactions are easy.
- Multi-shard transactions require two-phase commit or compensating transactions (sagas).

5. Failover & Replication

- Each shard should be replicated for high availability.
- Use leader-follower or multi-master setups depending on your consistency needs.

Real-World Examples & Tools

- **MySQL / PostgreSQL**: Proxy-based sharding with Vitess (MySQL) or Citus (Postgres).
- **MongoDB**: Built-in sharding using range or hashed shard keys and its balancer.
- **Cassandra**: Automatically partitions via consistent hashing—no manual shard management.
- **DynamoDB**: Client-chosen partition key with automatic rebalancing and replication.

When to Shard?

- Your dataset or traffic outgrows a single machine's CPU, RAM, or I/O.
- You need to isolate "hot" tenants or user ranges to avoid noisy neighbors.
- You require geo-distribution (shards in different regions).

TL;DR:

- **Partitioning** breaks data into chunks; **sharding** places each chunk on its own server.
- Choose a shard key that spreads load evenly.
- Plan for rebalancing, cross-shard operations, and replication to maintain performance and availability.

17. **Performance vs scalability**: A service is scalable if it results in increased performance in a manner proportional to resources added.

If you have a performance problem, your system is slow for a single user.

If you have a scalability problem, your system is fast for a single user but slow under heavy load.

18. **Consistent Hashing**

Consistent Hashing is a key technique used in distributed systems to distribute data across multiple nodes (servers, caches, etc.) in a way that **minimizes data movement** when nodes are added or removed.

What Problem Does It Solve?

Suppose you're using a basic hash function like:

```
server = hash(key) % N;
```

- If **N** (number of servers) changes, almost **all keys get remapped**, causing huge cache/database invalidations or data shuffles.
- That's not scalable for dynamic environments.

Consistent Hashing: The Idea

- Imagine a **circle (ring)** from **0** to **$2^{32} - 1$** .

- Both **keys and servers** are hashed to positions on this ring.
- A key is assigned to the **first server clockwise** from its position.

If a node is removed or added:

- Only a **small subset** of keys needs to move (those mapped to that node or near it).
-

Example

Suppose you have 4 servers S_1, S_2, S_3, S_4 placed on a hash ring.

A key like $user123$ hashes to a point on the ring:

- It will be assigned to the **next server clockwise**.

When you add S_5 :

- Only the keys that now fall between S_4 and S_5 move to S_5 .
 - Other keys stay with their current servers.
-

Virtual Nodes (Replicas)

To handle uneven load:

- Each server can be represented by **multiple points (virtual nodes)** on the ring.
 - This improves load balancing.
-

Benefits

- **Scalable:** Easy to add/remove nodes with minimal disruption.
 - **Fault-tolerant:** Fewer keys are affected if a node goes down.
 - **Widely used** in systems like:
 - Amazon DynamoDB
 - Apache Cassandra
 - Memcached
 - Kafka partitions
-

Time/Space Complexity

| Operation | Complexity |
|------------|---------------------------|
| Add node | $O(K/N)$ key remapping |
| Lookup key | $O(\log N)$ using TreeMap |

Insert/Delete

O(log N) with virtual nodes

19. Monolithic And Microservice Architecture

Monolithic Architecture

"All-in-one" application where all modules are part of a single codebase and deployable unit.

Pros:

- Easier to **develop, test, and deploy** initially.
- **Simple debugging** (single app logs, traces).
- Efficient **intra-process communication** (function calls).

Cons:

- Difficult to **scale specific parts** (must scale entire app).
- **Tightly coupled**: one change may affect everything.
- **Hard to maintain** as the codebase grows.
- Longer **deployment cycles** (especially for big teams).
- Technology stack is **locked** (e.g., can't mix Java and Node.js).

Microservices Architecture

"Divide and conquer" — application is broken into multiple small, independent services.

Each service:

- Owns its **logic and data**.
- Communicates over network (usually via REST, gRPC, etc.).
- Can be **developed, deployed, and scaled independently**.

Pros:

- **Scalability**: scale hot services independently.
- **Faster deployments** with CI/CD.
- **Tech flexibility**: different services can use different stacks.

- Better **fault isolation** — one failure doesn't kill the whole app.
- Teams can **own services independently**.

Cons:

- **Complexity**: network communication, API contracts, service discovery.
- **Data consistency** is harder.
- **Testing and debugging** are more complicated (distributed logs, tracing).
- **Deployment overhead** (multiple services = more infrastructure setup).

When to Use?

| Situation | Prefer |
|----------------------------|--|
| Small team/startup |  Monolith |
| Large-scale systems |  Microservices |
| Rapid prototyping |  Monolith |
| Need tech flexibility |  Microservices |
| Fast scaling of components |  Microservices |
| Simple CRUD app |  Monolith |

20. [Optimistic vs Pessimistic Locking](#)

Optimistic Locking

Assumes multiple transactions can complete without affecting each other. Conflict is checked only at commit time.

How it works:

- No actual lock is acquired.
- Each row has a version (e.g., timestamp or counter).
- On update, version is compared:
 - If version matches → update.
 - Else → **conflict detected**, retry or abort.

Pros:

- Great for **read-heavy systems**.
- **High performance** due to no blocking.

- Suitable for **distributed systems** and **low-contention scenarios**.

Cons:

- Conflict resolution may be **complex**.
- More suitable when **conflicts are rare**.

Example:

```
if (currentVersion == expectedVersion) {
    updateRecord(newData, currentVersion + 1);
} else {
    throw ConflictException;
}
```

Pessimistic Locking

Assumes conflicts are likely, so it locks resources upfront to prevent them.

How it works:

- Lock is acquired before accessing data.
- Other transactions are blocked until the lock is released.

Pros:

- **No risk of conflicts** during transaction.
- Ensures **strong consistency**.

Cons:

- Can lead to **deadlocks**.
- Poor performance in **high-concurrency** environments.
- **Locks** increase **latency** and block others.

Example:

```
-- Lock row for update
SELECT * FROM accounts WHERE id = 1 FOR UPDATE;
```

When to Use:

| Use Case | Strategy |
|------------------------------------|---|
| High read-to-write ratio |  Optimistic |
| High contention for data |  Pessimistic |
| Distributed system / stateless API |  Optimistic |
| Critical financial transactions |  Pessimistic |

21.  Strong vs Eventual Consistency

Strong Consistency

Guarantees that after an update, any read will return the most recent write.

Key Characteristics:

- **Linearizability:** All nodes agree on the same state at any given time.
- Operations appear **instantaneous and atomic**.
- **Read after write** is guaranteed to reflect latest data.

Use Cases:

- Banking systems 
- Inventory updates 
- Financial transactions or anything requiring absolute accuracy

Trade-offs:

- **Slower** due to coordination between nodes
- Reduced **availability** if quorum or leader node is unavailable (CAP: CA system)

Eventual Consistency

Guarantees that if no new updates are made, all replicas will eventually become consistent.

Key Characteristics:

- Reads may return **stale data**, but will converge eventually.
- **High availability and low latency**.

- Favored in **AP systems** in CAP theorem.

Use Cases:

- Social media timelines 
- DNS records 
- Shopping cart views 

Trade-offs:

- **Inconsistent reads** in short-term
- Requires **conflict resolution** logic sometimes

Summary Table:

| Feature | Strong Consistency  | Eventual Consistency  |
|----------------|--|--|
| Read freshness | Always up-to-date | May be stale |
| Availability | Lower (requires quorum) | Higher |
| Latency | Higher | Lower |
| Complexity | Lower | Higher (needs conflict resolution) |
| Use Cases | Finance, orders | Social feeds, cache |

22. Relational DB vs NoSQL

Relational Databases (RDBMS)

Structured data stored in tables with rows and columns. Uses SQL (Structured Query Language) for queries.

Characteristics:

- **Schema-based:** Fixed schema, predefined structure.
- **ACID compliant:** Atomicity, Consistency, Isolation, Durability.
- Uses **Joins** for relationships.
- Strong **data integrity**.

Examples:

- MySQL
- PostgreSQL

- Oracle
- SQL Server

Best For:

- Complex queries
- Structured data
- Banking, ERPs, inventory systems

NoSQL Databases

Designed for unstructured or semi-structured data. Flexible schema or schema-less.

Types:

1. **Document-based** (e.g., MongoDB)
2. **Key-Value stores** (e.g., Redis, DynamoDB)
3. **Column-family** (e.g., Cassandra, HBase)
4. **Graph** (e.g., Neo4j)

Characteristics:

- **Schema-less** or dynamic schema
- **Horizontal scaling** (easier to scale)
- **Eventually consistent** (CAP: AP or CP)
- High throughput for **large-scale data**

Examples:

- MongoDB
- Cassandra
- Redis
- Couchbase
- DynamoDB

Best For:

- Real-time analytics

- IoT and mobile apps
 - Social networks, chat applications
-



Summary Table:

| Feature | Relational DB (RDBMS) | NoSQL |
|----------------|-----------------------------|---------------------------------|
| Data Structure | Tables (rows, columns) | Documents, key-value, etc. |
| Schema | Fixed | Flexible / Dynamic |
| Query Language | SQL | Varies (e.g., JSON, query APIs) |
| Transactions | Full ACID | Limited or BASE |
| Scaling | Vertical | Horizontal |
| Best For | Structured, relational data | Unstructured, high-volume data |
| Joins | Yes | Not natively supported |



When to use what?

- Use **RDBMS** if:
 - Your data is **highly structured and relational**.
 - You need **strong consistency and ACID** guarantees.
- Use **NoSQL** if:
 - Your data is **unstructured or rapidly changing**.
 - You need **scalability**, high write/read throughput, or **flexibility**.

23. Types of NoSQL

- Document Based
- Key Value
- Wide Column
- Graph Based

◆ 1. Document-Based NoSQL Databases

- **Stores data as documents** (typically JSON, BSON, or XML).
- Each document is **self-describing**, flexible, and can have a unique structure.



Use Cases:

- Content management systems

- User profiles
- Catalogs

Examples:

- MongoDB
 - Couchbase
 - Firebase Firestore
-

◆ 2. Key-Value Stores

- Data is stored as **key-value pairs**.
- Great for **fast lookups**, session storage, and caching.

Use Cases:

- Caching (e.g., user session data)
- Shopping cart data
- Real-time recommendations

Examples:

- Redis
 - DynamoDB
 - Riak
 - Memcached
-

◆ 3. Column-Family Stores (Wide Column Stores)

- Stores data in **columns instead of rows**.
- Optimized for **read/write on large datasets**.
- Flexible schema for each row.

Use Cases:

- Time-series data
- Logging and analytics
- High-write systems

Examples:

- **Apache Cassandra**
 - HBase
 - ScyllaDB
-

◆ 4. Graph Databases

- Represents data as **nodes** (entities) and **edges** (relationships).
- Excellent for **relationship-heavy queries**.

✓ Use Cases:

- Social networks
- Fraud detection
- Recommendation engines

📌 Examples:

- **Neo4j**
 - ArangoDB
 - JanusGraph
-

Summary Table:

| NoSQL Type | Data Format | Strengths | Example |
|-----------------|-----------------|--------------------------------|-----------------|
| Document Store | JSON, BSON | Flexible schema, nesting | MongoDB |
| Key-Value Store | Key ↔ Value | Speed, simplicity | Redis, DynamoDB |
| Column Store | Column families | Scalability, large data writes | Cassandra |
| Graph Database | Nodes & Edges | Relationships, traversals | Neo4j |

24. Caching

Caching is a strategy used to store **frequently accessed data** in a temporary storage location (called a **cache**) to reduce **latency** and **improve performance**.

1. **Cache-aside**
2. **Write-through**
3. **Write-behind (write-back)**
4. **Refresh-ahead**

Why Use Caching?

-  **Reduce latency:** Faster access to data
 -  **Reduce load:** Fewer requests to backend/database
 -  **Cost-effective:** Less computation or external API calls
 -  **Boost scalability:** Improves system throughput
-

Where Can You Apply Caching?

| Layer | Example |
|----------------|---|
| Client-side | Browser localStorage, Service Workers |
| CDN/Edge | Cloudflare, Akamai (for static content) |
| API Layer | HTTP cache headers, API Gateway |
| App Layer | Method-level caching (e.g., Spring Cache) |
| Database Layer | Redis/Memcached for query results |
| Filesystem | OS-level disk cache, in-memory cache |

Popular Caching Tools

| Tool | Type | Use Case |
|-----------|----------------------|--------------------------------------|
| Redis | In-memory, key-value | Session, tokens, query results |
| Memcached | In-memory, key-value | Lightweight caching, flat data |
| CDNs | Edge caching | Static files (images, JS, CSS, etc.) |
| Ehcache | Java in-process | Local method-level or object caching |

Cache Expiry Strategies

- **TTL (Time to Live)** – Auto-expire after  seconds.
 - **LRU (Least Recently Used)** – Remove least-used item when full.
 - **LFU (Least Frequently Used)** – Remove least-frequent access items.
 - **FIFO (First In First Out)** – Remove oldest item.
-

Cache Invalidation Strategies

| Strategy | Description |
|---------------|---------------------------------------|
| Write-through | Data written to cache and DB together |

| | |
|------------------------------|---|
| Write-around | Only write to DB; cache updated on next read |
| Write-back (lazy) | Write to cache, then to DB after delay |
| Explicit invalidation | Manually delete or update cache on write/update |

⚠ Cache Challenges

- ✗ **Stale data** (if not invalidated properly)
- ⚠ **Cache stampede** (many requests on cache miss)
- 🚧 **Cold start** (initial cache has no data)
- 📊 **Consistency** with database

✓ Best Practices

- Use caching for **read-heavy workloads**
- Cache **idempotent, non-sensitive**, and **frequently requested** data
- Avoid caching data that **changes frequently** or needs **real-time accuracy**
- Always define a **cache eviction** policy

25. [Cache Eviction](#)

Cache eviction is the process of removing data from the cache when it's full or when certain policies dictate it's no longer needed. Since memory (cache) is limited, we can't store everything forever — eviction ensures efficient use of that space.

🔄 Common Cache Eviction Policies

| Policy | Full Form | When it removes | Best Used When |
|---------------|-----------------------|---|---|
| LRU | Least Recently Used | Removes the item that hasn't been used for the longest time | Access patterns are temporal (recently accessed = more likely to be reused) |
| LFU | Least Frequently Used | Removes item with the lowest access frequency | Some data is accessed frequently, rest rarely |
| FIFO | First In First Out | Removes the oldest added item | No priority or frequency requirement |
| MRU | Most Recently Used | Removes the most recently accessed item | Rarely used, good in specific cases (e.g. undo buffers) |
| Random | — | Removes a random item | Simple, fast, less overhead; useful in some distributed caches |

Eviction in Popular Tools

| Tool | Default Eviction | Notes |
|---------------|------------------------------------|--|
| Redis | Configurable: LRU, LFU, FIFO, etc. | <code>maxmemory-policy</code> in <code>redis.conf</code> |
| Memcached | LRU | Simple, non-persistent |
| Browser Cache | LRU + heuristics | Based on resource type, size, and priority |

Extra Concepts

- **Write Policies:**
 - *Write-Through*: Updates cache and DB simultaneously
 - *Write-Back*: Updates cache first, then DB later (less safe)
 - *Write-Around*: Skips cache on write, writes only to DB
- **Cache Miss Penalty**: If the cache eviction removes a useful item and it's needed again, a DB call (which is slower) is made — leading to latency.
- **Multi-Level Caching**: L1 (in-memory) and L2 (Redis/Distributed) — each level may have its own eviction logic.

25. `Data Center / Racks / Hosts`

1. Data Center

A **data center** is a **physical facility** used to house computing resources like servers, networking equipment, and storage systems. It typically contains:

- Power and cooling systems
- Physical security
- Firewalls and switches
- Backup generators
- Redundant internet and power lines

Think of it as a large warehouse full of computing power.

2. Rack

Inside a data center, servers are **organized into vertical units called "racks."**

- A **rack** is a **frame** that holds multiple **physical servers** (hosts).
- Each rack typically has:

- 42U (standard rack height — 1U = 1.75 inches)
- Network switch at the top (ToR switch)
- Power Distribution Units (PDUs)

| Think of a rack as a bookshelf, and servers as books inside it.

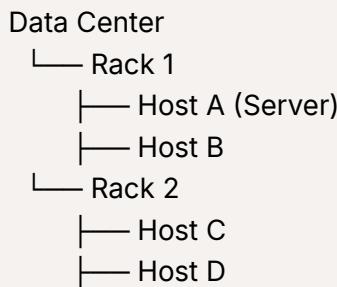
3. Host

A **host** (also known as a **node** or **physical server**) is:

- A **machine** with CPU, memory, storage, and network capabilities
- Can run:
 - A single OS (bare-metal server)
 - Multiple Virtual Machines (via hypervisors)
 - Containers (via Docker, Kubernetes, etc.)

| Think of a host as the actual computer doing the work.

Hierarchy Summary



Cloud Analogy

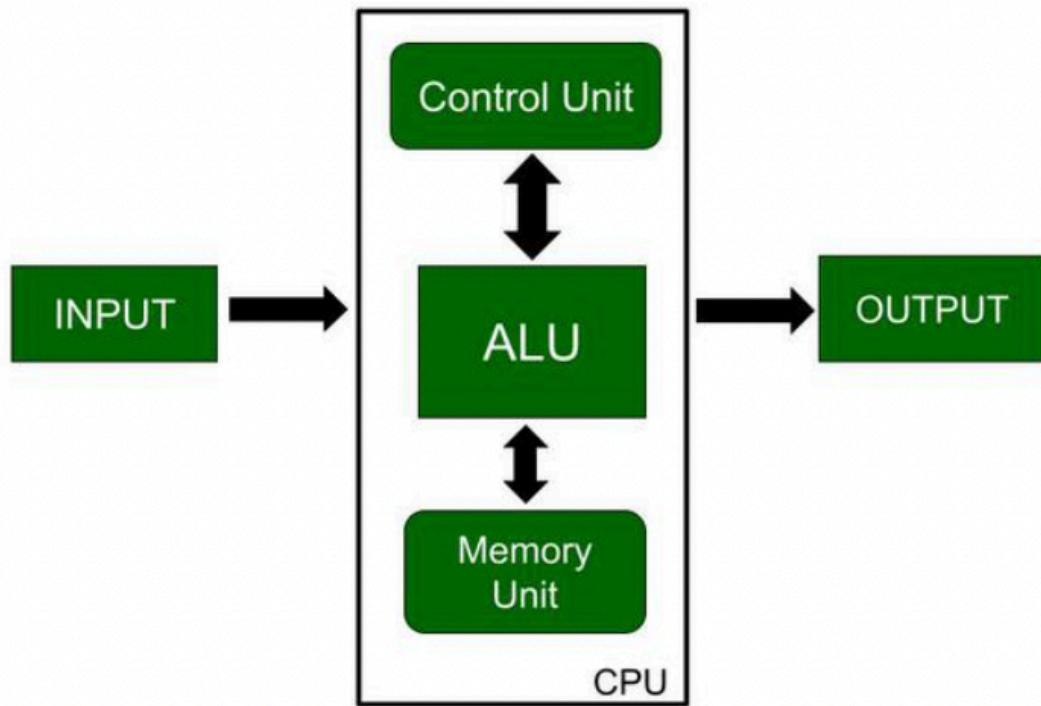
In cloud providers like AWS, GCP, Azure:

- **Availability Zones (AZs)** ≈ different data centers
- **Physical host** = the machine your VM runs on
- **Racks and hosts** are abstracted from users, but they still exist physically

26. CPU / Memory / Hard Drive / Network Bandwidth

1. CPU (Central Processing Unit)

- **Role:** Executes instructions (code).
- **Bottleneck Symptoms:** Slow computation, high latency, thread contention.
- **Optimization Tips:**
 - Use efficient algorithms.
 - Parallelize tasks (multithreading/multiprocessing).
 - Profile CPU usage using tools like `top`, `htop`, or `perf`.



2. Memory (RAM)

- **Role:** Stores data temporarily while a program is running.
- **Bottleneck Symptoms:** Out-of-Memory (OOM) errors, frequent swapping, crashes.
- **Optimization Tips:**
 - Use caching wisely.
 - Clean up unused objects.
 - Use memory profilers (`Valgrind`, `VisualVM`, etc.).

3. Hard Drive (Disk I/O)

- **Role:** Permanent storage for files, databases, logs.
 - **Bottleneck Symptoms:** Slow reads/writes, high I/O wait time.
 - **Optimization Tips:**
 - Use SSDs over HDDs.
 - Optimize DB queries.
 - Batch write operations.
-

4. Network Bandwidth

- **Role:** Transfers data between systems or services.
 - **Bottleneck Symptoms:** Slow API responses, packet drops, timeouts.
 - **Optimization Tips:**
 - Compress payloads (GZIP).
 - Use CDN for static assets.
 - Optimize APIs and reduce redundant requests.
-

Summary Table:

| Resource | Primary Use | Bottleneck Symptom | Optimization Idea |
|------------|---------------------------|----------------------|------------------------------------|
| CPU | Executes instructions | High CPU, slow logic | Multithreading, profiling |
| Memory | Temporary data storage | OOM, crashes | Efficient data structures, caching |
| Hard Drive | Persistent data storage | Slow I/O, timeouts | SSD, DB query optimization |
| Network | Remote data communication | Lag, dropped packets | Compression, pagination, CDNs |

27. Random vs Sequential Read/Write on Disk

Sequential Read/Write

- **Definition:** Data is read from or written to **contiguous** disk blocks.
- **Performance:** 🔥 **Faster** (especially on HDDs)
- **Reason:** Minimal disk head movement; OS and disk can optimize for prefetching.
- **Use Cases:**

- Log files
 - Video/audio streaming
 - Backup dumps
 - Append-only databases
-

Random Read/Write

- **Definition:** Data is read/written at **non-contiguous** positions on disk.
 - **Performance:**  **Slower** (especially on HDDs)
 - **Reason:** Requires frequent seeks and disk head repositioning.
 - **Use Cases:**
 - Database queries (random access to records)
 - File systems with frequent updates
 - Index access in databases
-

SSD vs HDD:

| Operation Type | HDD Speed | SSD Speed |
|------------------|--|---|
| Sequential Read |  Fast (~150 MB/s) |   Faster (~500+ MB/s) |
| Random Read |  Slow (~1 IOPS) |   Fast (100K+ IOPS) |
| Sequential Write |  Fast |   Faster |
| Random Write |  Slow |   Fast |

SSDs handle random I/O much better due to no mechanical parts.

Tip for System Design:

- Design **write-heavy** systems (like logs, metrics) to use **sequential writes**.
- Use **B-Trees** or **LSM-Trees** depending on random vs sequential patterns in DBs.
- Prefer SSDs for **high-performance, random-access workloads**.

28. [HTTP vs HTTP2 vs WebSockets](#)

HTTP (1.1) – The Traditional Web Protocol

- **Request/Response Model:** Client requests, server responds.
- **Connection:** One request per TCP connection (or pipelined, but rarely used effectively).

- **Latency:** High due to multiple round trips.
 - **Headers:** Sent in plain text (repeated headers, no compression).
 - **Use Case:** Simple web pages, REST APIs.
-

HTTP/2 – Faster, Efficient

- **Multiplexing:** Multiple streams over one TCP connection (parallel requests/responses).
 - **Header Compression:** Using HPACK → reduces overhead.
 - **Binary Protocol:** Unlike HTTP/1.1 which is text-based.
 - **Server Push:** Server can push resources before client asks.
 - **Use Case:** Modern websites, microservices (gRPC uses HTTP/2), streaming REST APIs.
-

WebSockets – Full-Duplex Communication

- **Persistent Connection:** After initial HTTP handshake, it upgrades to WebSocket.
 - **Full Duplex:** Client and server can send messages independently and simultaneously.
 - **Low Latency:** No need to establish a new connection for each interaction.
 - **Use Case:** Real-time apps (chat, live notifications, multiplayer games, stock tickers).
-

TL;DR Comparison Table:

| Feature | HTTP/1.1 | HTTP/2 | WebSockets |
|--------------------|------------------------|----------------------|-------------------------|
| Protocol Type | Text-based | Binary | Binary (after upgrade) |
| Communication | Half-duplex (req-resp) | Multiplexed req-resp | Full duplex |
| Connections Needed | Multiple | Single | Single (persistent) |
| Latency | Higher | Lower | Lowest |
| Best For | Websites, REST APIs | Modern apps, gRPC | Real-time communication |

When to use what?

- **HTTP/1.1:** Legacy systems, simple APIs.
- **HTTP/2:** For performance in modern web apps, gRPC-based services.
- **WebSockets:** For real-time and event-driven communication (chat, gaming, dashboards).

The **TCP/IP Model** (also known as the Internet Protocol Suite) defines how data is transmitted across networks, like the internet. It's the backbone of all network communication.

📦 Layers of TCP/IP Model (4-Layer Architecture)

| Layer | Purpose | Protocols/Examples |
|-----------------------------------|---|---|
| 1. Application Layer | Interfaces directly with the application. Handles protocols used by software apps. | HTTP, FTP, SMTP, DNS, WebSockets |
| 2. Transport Layer | Ensures reliable or unreliable delivery. Manages end-to-end communication. | TCP (reliable), UDP (fast but unreliable) |
| 3. Internet Layer | Responsible for addressing and routing data. | IP (IPv4/IPv6), ICMP, ARP |
| 4. Link (Network Interface) Layer | Deals with the physical transmission of data. | Ethernet, Wi-Fi, ARP (also fits here) |

🔄 Data Flow (Simplified)

- **Sender side:** App Data → Application → Transport → Internet → Link → 
 - **Receiver side:**  → Link → Internet → Transport → Application → App Data
-

🧠 Mnemonic to Remember:

All Teachers Inspire Learning

(Application, Transport, Internet, Link)

💡 TCP/IP vs (Open Systems Interconnection) OSI

| TCP/IP Model | OSI Model (7 Layers) |
|--------------|------------------------------------|
| Application | Application, Presentation, Session |
| Transport | Transport |
| Internet | Network |
| Link | Data Link + Physical |

💡 Real-World Mapping:

- You visit a website →  uses **HTTP**
- **TCP** ensures your request arrives intact → 
- **IP** routes the packets to the right server → 
- **Ethernet/Wi-Fi** moves bits physically → 

30. [IPv4 vs IPv6](#)

What is IP?

IP (Internet Protocol) provides **unique addresses** to devices on a network so they can communicate.

IPv4 vs IPv6: Key Differences

| Feature | IPv4 | IPv6 |
|------------------------|---|--|
| Address Length | 32-bit | 128-bit |
| Address Format | Decimal (e.g., <code>192.168.0.1</code>) | Hexadecimal (e.g., <code>2001:0db8::1</code>) |
| Total Addresses | ~4.3 billion | ~340 undecillion (2^{128}) |
| Header Size | 20 bytes | 40 bytes |
| Security | Optional (IPSec is optional) | Mandatory (IPSec is built-in) |
| Configuration | Manual/DHCP | Auto-configuration (SLAAC) |
| Broadcast | Yes | No (uses multicast/anycast) |
| NAT Required | Yes (due to address exhaustion) | No (more addresses available) |
| Checksum Field | Yes | No (removed for performance) |
| Fragmentation | Done by sender & routers | Only by sender |
| Deployment | Widely used | Gradual adoption |

Examples:

- **IPv4:**

`192.168.1.1`

- **IPv6:**

`2001:0db8:85a3:0000:0000:8a2e:0370:7334`

(can be shortened using `::` as: `2001:db8:85a3::8a2e:370:7334`)

Why IPv6?

- IPv4 can't keep up with the **explosive growth of devices**.
- IPv6 simplifies routing, enhances security, and removes the need for NAT.

31. [TCP vs UDP](#)

TCP vs UDP: The Basics

| Feature | TCP (Transmission Control Protocol) | UDP (User Datagram Protocol) |
|--|---|------------------------------------|
| Connection Type | Connection-oriented (handshake) | Connectionless |
| Reliability | Reliable (guarantees delivery, order, no dupes) | Unreliable (no guarantees) |
| Ordering | Maintains order of packets | No order preservation |
| Speed | Slower (due to overhead) | Faster (minimal overhead) |
| Overhead | High (ack, sequence, retransmission, etc.) | Low |
| Error Checking | Yes (and retransmits if needed) | Yes (but no retransmission) |
| Use Cases | File transfer, email, web browsing (HTTP/HTTPS) | Video streaming, gaming, VoIP, DNS |
| Header Size | 20–60 bytes | 8 bytes |
| Flow Control / Congestion Control | Yes | No |
| Packet Acknowledgement | Yes | No |

TCP Lifecycle (3-Way Handshake):

1. SYN
 2. SYN-ACK
 3. ACK
- Connection established.

When to Use:

- **TCP**: When **reliability** is crucial (e.g. banking apps, HTTPS).
- **UDP**: When **speed** matters more than accuracy (e.g. live video, online games).

Memory Tip:

TCP = Take Care of Packets

UDP = Unreliable but Damn fast Protocol

UDP is a **connectionless, lightweight protocol** in the transport layer (TCP/IP model) used for **fast, real-time** communication where **low latency** is more important than reliability.

No handshake; data sent without setup

Unlike TCP, **UDP (User Datagram Protocol)** does **not establish a connection** before sending data. It's a **fire-and-forget** style protocol — meaning:

- Data is sent directly to the recipient
- No handshaking
- No connection state
- No acknowledgment or retransmission
- No built-in reliability

So how does "UDP communication" begin?

Even though UDP is connectionless, here's a typical **flow of communication**:

1. **Sender prepares a datagram**
 - Includes destination IP and port
2. **Sends data to receiver**
 - Via UDP socket (e.g., `sendto()` in C, or `socket.send()` in Node.js)
3. **Receiver listens on a port**
 - Via UDP socket (e.g., `recvfrom()`)
4. **Receiver gets data**
 - No acknowledgment sent back by default

Sending a UDP packet is like mailing a postcard with no guarantee it'll arrive — and no return address needed.

TCP: Connection-Oriented Protocol

TCP **establishes a reliable connection** before any data is exchanged. It ensures that data is delivered **accurately** and **in order**.

TCP Connection Establishment (3-Way Handshake)

This is how a TCP connection is formed:

1. SYN:

Client sends a SYN (synchronize) packet to the server to initiate a connection.

2. SYN-ACK:

Server responds with a SYN-ACK (synchronize-acknowledge) packet.

3. ACK:

Client replies with an ACK (acknowledge) packet.

- ✓ Connection is now established and both sides can begin communication.
-

Features of TCP:

-  **Reliable**: Guarantees delivery via acknowledgments and retransmissions.
 -  **Ordered**: Maintains sequence of packets.
 -  **Stateful**: Keeps track of connection state.
 -  **Flow Control**: Prevents sender from overwhelming receiver.
 -  **Congestion Control**: Adapts speed based on network conditions.
-

Analogy:

TCP is like a phone call. You greet, the other person answers, you start talking. You know if the message was received.

Use Cases:

- Web (HTTP, HTTPS)
 - File transfers (FTP, SFTP)
 - Emails (SMTP, IMAP, POP3)
 - Database communication (MySQL, PostgreSQL)
-

Downsides:

- Slower than UDP due to overhead
- Consumes more memory and CPU for tracking state

32. 

What is DNS Lookup?

DNS (Domain Name System) Lookup is the process of translating a **domain name** (like `google.com`) into an **IP address** (like `142.250.77.206`) that computers can use to locate each other on the internet.

Steps in a DNS Lookup:

1. Browser Cache

- Browser checks if domain is cached locally.

2. OS Cache

- If not found, the operating system checks its own DNS cache.

3. Router Cache

- If OS doesn't know, query goes to the router's DNS cache.

4. ISP's Recursive DNS Server

- Router sends query to ISP's recursive resolver.

5. Root DNS Server

- Resolver queries a **root server** → returns address of TLD server ([.com](#)).

6. TLD DNS Server

- Resolver queries **TLD server** → returns address of authoritative server ([google.com](#)).

7. Authoritative DNS Server

- Resolver queries **authoritative server** → gets final IP address.

8. IP Returned

- Resolver caches the result and sends IP to the browser → connection is established.

🧠 DNS Query Types:

| Type | Purpose |
|-------|----------------------------|
| A | IPv4 address |
| AAAA | IPv6 address |
| MX | Mail servers |
| CNAME | Canonical name (alias) |
| NS | Name servers |
| TXT | Metadata (e.g., SPF, DKIM) |

📌 DNS Lookup Example:

google.com → DNS → 142.250.77.206 → Browser connects

⚡ Tools to Test DNS Lookup:

- nslookup google.com
- dig google.com
- host google.com

33. [HTTPS & TLS](#)

HTTPS (HyperText Transfer Protocol Secure)

| It's just HTTP + TLS encryption = Secure HTTP.

TLS (Transport Layer Security)

TLS is a cryptographic protocol that:

-  **Encrypts** data between client and server.
 -  Ensures **data integrity** (not tampered with).
 -  **Authenticates** the server (and optionally the client).
-

TLS Handshake (Simplified)

1. Client Hello

- Client says "Hi, here are the TLS versions & ciphers I support."

2. Server Hello

- Server chooses cipher, sends its **digital certificate** (proof it's real).

3. Key Exchange

- Client and server exchange keys (usually using Diffie-Hellman or RSA).

4. Session Key Established

- Both now share a symmetric session key to encrypt all future data.

5. Encrypted Communication Starts

- Secure HTTPS communication begins.
-

TLS Features:

-  **Encryption:** All traffic is encrypted.
 -  **Authentication:** Confirms identity of the server via certificates.
 -  **Integrity:** Ensures the data wasn't changed en route.
-

🔑 TLS vs SSL

- SSL is old and insecure.
- TLS is the modern replacement (TLS 1.3 is the latest standard).

✓ Why HTTPS matters:

- Prevents **eavesdropping**, **man-in-the-middle attacks**, and **data leaks**.
- Required for **secure logins**, **credit card payments**, etc.

🧠 Analogy:

HTTP is like sending a postcard—anyone can read it.

HTTPS is like sending a locked box—only the receiver has the key.

🔒 SSL (Secure Sockets Layer)

SSL is a protocol designed to **secure communication over the internet** by **encrypting data** between clients (like browsers) and servers.

📌 SSL is deprecated. Modern systems use TLS (Transport Layer Security), but the term "SSL" is still commonly used.

💡 What SSL Does:

- Ensures **confidentiality**: Data is encrypted so no one else can read it.
- Provides **integrity**: Detects if the data was tampered with.
- Offers **authentication**: Verifies that you're talking to the legitimate server.

🔒 How SSL/TLS Works:

1. Handshake Phase:

- The client sends a `ClientHello` message with supported ciphers.
- The server replies with a `ServerHello`, picks a cipher, and sends its **SSL certificate**.
- The client validates the certificate (via CA).
- A **session key** is exchanged (via asymmetric encryption).

2. Session Phase:

- All further communication is encrypted using the shared session key (symmetric encryption).
-

Versions:

- SSL 1.0 – Never released.
 - SSL 2.0 – Obsolete.
 - SSL 3.0 – Deprecated.
 - **TLS 1.0 → 1.1 → 1.2 → 1.3** – Secure and widely used today.
-

Uses of SSL/TLS:

- HTTPS (Secure Web Browsing)
 - Secure Email (SMTPTS, IMAPS)
 - Secure FTP (FTPS)
 - VPNs
-

Key Terms:

- **SSL Certificate:** Contains public key + identity of server (issued by Certificate Authority)
- **CA (Certificate Authority):** Trusted entities that issue SSL certificates
- **HTTPS = HTTP + SSL/TLS**

35. Public Key Infrastructure & Certificate Authority

Public Key Infrastructure (PKI)

PKI is a framework that manages digital keys and certificates to facilitate secure communication, authentication, and encryption over a network. It relies on both **public** and **private** keys to ensure the integrity and security of data exchange.

Core Components of PKI:

1. **Public and Private Keys:**
 - **Public Key:** Available for anyone to encrypt data or verify signatures.
 - **Private Key:** Kept secret by the owner, used to decrypt data or create digital signatures.
2. **Digital Certificates:**
 - A digital certificate binds the public key to an entity (like an individual, organization, or server).
 - It includes:

- Public key
- Information about the certificate holder
- Certificate Authority (CA) details
- Validity period
- Digital signature from the CA

3. Certificate Authority (CA):

- A trusted organization that issues digital certificates.
- The CA verifies the identity of the certificate requestor before issuing a certificate.
- CAs sign certificates with their own private key, and these signed certificates can then be validated by anyone using the CA's public key.

4. Registration Authority (RA):

- The RA is responsible for accepting requests for digital certificates and authenticating the entity making the request.
- It acts as a mediator between the user and the CA.

5. Digital Signature:

- A cryptographic technique that verifies the authenticity and integrity of a message, document, or transaction.
- It ensures that the sender's identity is confirmed and that the message has not been altered.

6. Certificate Revocation List (CRL):

- A list of certificates that have been revoked by the CA before their expiration date.
- CRLs are used to check whether a certificate is still valid.

7. Key Pair Management:

- The secure generation, storage, and management of public and private keys.
- Typically involves using hardware or software tools like **HSMs (Hardware Security Modules)**.

How PKI Works (In Practice):

1. Key Generation:

- A user generates a pair of keys: one public and one private.

2. Certificate Request:

- The user sends a **Certificate Signing Request (CSR)** to a CA, including their public key and identity information.

3. Certificate Issuance:

- After the CA authenticates the user's identity, it issues a digital certificate, containing the public key and the CA's signature.

4. Encryption & Decryption:

- The public key can be used to encrypt data, which can only be decrypted by the corresponding private key.
- Similarly, the private key can be used to create a digital signature, which others can verify using the public key.

5. Authentication:

- When the certificate is used, the recipient can verify the signature using the public key contained in the certificate to ensure the message was sent by the rightful sender.

Certificate Authority (CA)

A **Certificate Authority (CA)** is a trusted third party in the PKI that issues and manages digital certificates. The CA ensures the identity of the certificate holder by performing a rigorous validation process.

Role of a CA:

1. Issuing Certificates:

- The CA signs certificates after validating the identity of an entity.

2. Revoking Certificates:

- If a certificate is compromised or no longer valid, the CA revokes it and publishes a **Certificate Revocation List (CRL)**.

3. Certificate Renewal:

- CAs manage the renewal process for certificates, ensuring they stay valid.

4. Root and Intermediate CAs:

- **Root CA:** The top-level CA that is trusted by browsers and devices. Root certificates are pre-installed on devices and browsers.
- **Intermediate CA:** Issues certificates on behalf of the root CA. It acts as a bridge between the root CA and the end-user certificate.

Public Key Infrastructure in Action (SSL/TLS Example):

1. A **client** (e.g., a web browser) connects to a **server** (e.g., a website).

2. The **server** sends its **digital certificate** containing its public key.
 3. The **client** verifies the certificate by checking if it's signed by a trusted CA.
 4. If the certificate is valid, the **client** uses the server's public key to encrypt data (e.g., session keys).
 5. The **server** decrypts the data with its private key, and the communication proceeds securely.
-



Types of CAs:

1. **Root CA:**
 - The highest authority in the PKI hierarchy. Its public key is self-signed and trusted by default.
 2. **Intermediate CA:**
 - A CA that is signed by the root CA and issues certificates to end users.
 3. **Public CA:**
 - A CA that issues certificates for public use. Examples include Let's Encrypt, DigiCert, and GlobalSign.
 4. **Private CA:**
 - A CA that issues certificates for private or internal use, typically within an organization.
-



Benefits of PKI and CA:

- **Confidentiality:** Ensures that only authorized parties can access sensitive data.
 - **Integrity:** Guarantees that the data hasn't been altered.
 - **Authentication:** Confirms the identity of communicating parties.
 - **Non-repudiation:** Ensures that the sender cannot deny having sent the message.
-

36. Symmetric vs Asymmetric Key



Symmetric vs Asymmetric Key Cryptography

Cryptography is the practice of securing communication and data through the use of codes. Symmetric and asymmetric key cryptography are two of the main methods used to secure data. Here's a breakdown of both:

1. Symmetric Key Cryptography (Secret Key Cryptography)

In **symmetric cryptography**, the same key is used for both **encryption** and **decryption** of data. The key must be kept secret and shared between the sender and the receiver.

How It Works:

- The sender uses the **shared secret key** to encrypt the message.
- The receiver uses the **same shared key** to decrypt the message.

Characteristics:

- **Speed:** Symmetric encryption algorithms are generally faster than asymmetric algorithms.
- **Security Risk:** The biggest challenge is the secure distribution of the secret key. If the key is intercepted during transmission, the security is compromised.
- **Common Algorithms:** AES (Advanced Encryption Standard), DES (Data Encryption Standard), 3DES (Triple DES), RC4, Blowfish.

Use Cases:

- **Data-at-Rest:** Protecting data stored on disks.
- **Data-in-Transit:** Encrypting data over secure channels like HTTPS.
- **Bulk Encryption:** Symmetric encryption is often used to encrypt large amounts of data efficiently, once a secure key exchange has occurred.

Example:

- **AES (Advanced Encryption Standard):** One of the most widely used symmetric encryption algorithms, providing secure data encryption for various applications.

2. Asymmetric Key Cryptography (Public-Key Cryptography)

In **asymmetric cryptography**, a **pair of keys** is used: one **public key** and one **private key**. The keys are mathematically related but it is not possible to derive the private key from the public key.

How It Works:

- The sender uses the **receiver's public key** to encrypt the message.
- The receiver uses their **private key** to decrypt the message.

Characteristics:

- **Two Keys:** There is a pair of keys involved — **public key** (used for encryption) and **private key** (used for decryption).

- **Security:** Public keys can be distributed freely, and private keys are kept secret. This makes it more secure for data transmission, as only the holder of the private key can decrypt the message.
- **Speed:** Asymmetric encryption is slower than symmetric encryption, because the mathematical operations involved are more computationally intensive.
- **Key Distribution:** The major advantage over symmetric encryption is that you don't need to securely distribute the private key, only the public key.



Use Cases:

- **Digital Signatures:** Verifying the authenticity of a message or document (using private key for signing and public key for verification).
- **Key Exchange:** Using asymmetric cryptography (e.g., RSA) to securely exchange symmetric keys over an insecure channel (e.g., SSL/TLS).
- **Email Encryption:** Ensuring secure communication via encrypted email.

Common Algorithms:

- **RSA:** One of the most widely used asymmetric encryption algorithms.
- **ECC (Elliptic Curve Cryptography):** A newer asymmetric encryption technique that is more efficient and secure than RSA for shorter key sizes.
- **DSA (Digital Signature Algorithm):** Mainly used for digital signatures.

Example:

- **RSA (Rivest-Shamir-Adleman):** A widely used asymmetric encryption algorithm used in securing data during transmission, such as in HTTPS.

Comparison: Symmetric vs Asymmetric Key Cryptography

| Feature | Symmetric Key Cryptography | Asymmetric Key Cryptography |
|--------------------------|--|--|
| Key Used | Same key for both encryption and decryption | A pair of keys (public & private) |
| Speed | Fast | Slower, due to more complex math |
| Security Risk | Risk of key interception during transmission | Public key is freely distributed, private key is kept secure |
| Key Distribution | Secure key distribution required | Public key can be freely shared, private key remains secret |
| Common Algorithms | AES, DES, 3DES, RC4, Blowfish | RSA, ECC, DSA, ElGamal |

| | | |
|-------------------|--|---|
| Use Cases | Data-at-rest, secure communications | Digital signatures, email encryption, key exchange (e.g., in SSL/TLS) |
| Efficiency | More efficient for large amounts of data | More efficient for key exchange but less for bulk data encryption |

When to Use Symmetric vs Asymmetric?

- **Symmetric encryption** is typically used for **encrypting large volumes of data** quickly once the **key exchange** has been securely performed (for example, during a **TLS handshake**).
- **Asymmetric encryption** is commonly used for **key exchange** (e.g., in **SSL/TLS protocols**) or for **digital signatures**, where one key is used to verify the identity of the sender.

Hybrid Approach (Common in SSL/TLS)

In many modern systems (such as SSL/TLS for securing web traffic), both symmetric and asymmetric cryptography are used together:

1. **Asymmetric cryptography** is used in the handshake phase to **securely exchange a symmetric key**.
2. **Symmetric encryption** is then used to **encrypt the actual data** for performance reasons (since symmetric encryption is faster).

This hybrid approach combines the best of both worlds: the secure key distribution of asymmetric cryptography and the speed of symmetric encryption.

37.  Load Balancer → L4 vs L7

A **Load Balancer** distributes incoming traffic across multiple servers to ensure no single server becomes a bottleneck. The two common types of load balancers are:

Layer 4 Load Balancer (L4)

Layer 4 operates at the **Transport Layer** (TCP/UDP) of the OSI model.

How it Works:

- Makes routing decisions based on **IP address, TCP/UDP ports**.
- It doesn't look into the actual contents of the packets (no awareness of HTTP headers, cookies, etc.).

Advantages:

- **Faster and lightweight**: Works directly with network packets.
- **Protocol-agnostic**: Works with any TCP/UDP traffic.

- Lower latency and higher throughput.

Limitations:

- No content-based routing (e.g., can't route based on `/api` or `/images`).
- No SSL termination or smart features like cookie-based affinity.

Examples:

- AWS Network Load Balancer (NLB)
- HAProxy (Layer 4 mode)
- Linux IPVS

◆ Layer 7 Load Balancer (L7)

Layer 7 operates at the **Application Layer** (HTTP/HTTPS).

How it Works:

- Makes routing decisions based on **URL, headers, cookies, content type**, etc.
- Can perform actions like **SSL termination, compression, caching, and WAF (Web Application Firewall)**.

Advantages:

- Content-aware routing (e.g., `/api` → backend A, `/images` → backend B).
- Smart features like:
 - Redirects
 - URL rewrites
 - Cookie/session-based routing
- Can terminate SSL and forward unencrypted traffic internally (optional).

Limitations:

- Slower than L4 due to deep packet inspection.
- Limited to application-layer protocols (mainly HTTP/S, WebSockets, gRPC).

Examples:

- AWS Application Load Balancer (ALB)
- NGINX

- Traefik
 - Envoy
-

Quick Comparison:

| Feature | L4 Load Balancer | L7 Load Balancer |
|-----------------------|----------------------------|-----------------------------------|
| OSI Layer | Transport Layer (Layer 4) | Application Layer (Layer 7) |
| Protocol Support | TCP, UDP | HTTP, HTTPS, gRPC, WebSockets |
| Routing Criteria | IP, Port | URL path, headers, cookies |
| Performance | Faster | Slightly slower (deep inspection) |
| SSL Termination | ✗ Not supported (mostly) | ✓ Supported |
| Cookie-based Sessions | ✗ No | ✓ Yes |
| Use Case Examples | Generic traffic (DB, SMTP) | Web traffic, APIs, microservices |

When to Use What?

- Use **L4** when:
 - You need **high performance** and low latency.
 - You're dealing with non-HTTP traffic (TCP/UDP based protocols).
 - You want basic IP+port-based routing.
- Use **L7** when:
 - You want **smart routing** based on request data.
 - You need features like **SSL termination**, **WAF**, or **session persistence**.
 - You're working with modern **microservices** or **APIs**.

1. Layer 4 Load Balancer (L4) – Transport Layer

- Works at the **TCP/UDP** level
- Routes traffic based on:
 - IP address
 - Port number
- Doesn't inspect request content

Example:

- A client hits `LB_IP:80` → traffic is forwarded to `AppServer1:80`

Tools:

- HAProxy (L4 mode)
 - AWS Network Load Balancer
 - IPVS (IP Virtual Server)
-

2. Layer 7 Load Balancer (L7) – Application Layer

- Works at the **HTTP/HTTPS** level
- Routes based on:
 - URL paths (`/api/v1/user`)
 - Cookies
 - Headers
 - Hostnames
- Supports advanced features like:
 - SSL termination
 - Content-based routing
 - Compression, caching, rewriting

Example:

- `/api/*` → API server
- `/static/*` → Static file server

Tools:

- NGINX
 - Traefik
 - AWS Application Load Balancer (ALB)
 - Envoy
-

3. Global Load Balancer / Geo Load Balancer

- Routes traffic across **data centers/regions**
- Based on **geolocation, latency, health, or load**

Examples:

- Cloudflare Load Balancer
- AWS Route 53 (latency routing)

- GCP Cloud Load Balancer (global)
-

4. Hardware Load Balancer

- Proprietary physical appliances
- High throughput, used in traditional data centers

Examples:

- F5 Big-IP
 - Citrix NetScaler
-

Summary Table

| Type | Layer | Routing Criteria | Common Tools |
|--------------|-------------|---------------------------|----------------------|
| L4 | Transport | IP, Port | HAProxy, NLB |
| L7 | Application | Path, Header, Host | NGINX, ALB, Traefik |
| Geo / Global | DNS Level | Location, latency | Route 53, Cloudflare |
| Hardware | Physical | High-speed packet routing | F5, NetScaler |

38. CDNs & Edge

What is a CDN (Content Delivery Network)?

A **CDN** is a distributed network of servers placed in **various geographic locations** that **cache and serve content** (like images, videos, HTML, CSS, JS) to users **from the nearest server**, improving **load times and reliability**.

Benefits:

- **Faster content delivery** (reduced latency)
- **Offloads traffic** from your origin server
- **Reduces bandwidth costs**
- **Protects against DDoS attacks** (some CDNs offer security)
- **Improves SEO and user experience**

How it works:

1. User visits your site.
2. DNS resolves to the **nearest CDN edge server**.
3. If the content is cached there → served directly.

4. If not, it fetches from the origin, caches it, and serves it.

Examples:

- Cloudflare
- Akamai
- AWS CloudFront
- Google Cloud CDN
- Fastly

What is Edge Computing?

Edge Computing means **processing data closer to the user/device** instead of relying on a centralized cloud/data center. It's ideal for **real-time**, **latency-sensitive**, or **offline-first** apps.

Think: "Do as much work as possible close to where the data is generated."

Use Cases:

- IoT devices
- Smart cars
- Industrial sensors
- AR/VR apps
- Live video streaming
- Real-time analytics

How it's different from CDN:

| Feature | CDN | Edge Computing |
|--------------|------------------------------------|---|
| Purpose | Content caching & delivery | Data processing near the source |
| Focus | Speeding up delivery | Reducing processing latency |
| Primary Role | Serve static assets faster | Run code or services close to users |
| Examples | Images, CSS, JS files | Real-time analytics, ML inference, IoT |
| Tools | Cloudflare CDN, Akamai, CloudFront | Cloudflare Workers, AWS Lambda@Edge, Fly.io |

Together: CDN + Edge = Superpower

Modern CDNs like **Cloudflare**, **Fastly**, **Netlify**, **Vercel** also offer **Edge Functions** or **Edge Workers**, letting you:

- Do **A/B testing** at the edge
- **Personalize** content without hitting backend
- Handle **authentication, redirects, and rewrites**
- Serve **dynamic content** with very low latency

💡 TL;DR

| | | |
|----------|------------------------------------|---|
| Term | CDN | Edge Computing |
| Goal | Speed up content delivery | Run code close to user/device |
| Type | Mainly caching & distribution | Mainly logic & processing |
| Latency | Reduces content delivery latency | Reduces computation response latency |
| Examples | CloudFront, Akamai, Cloudflare CDN | Lambda@Edge, Cloudflare Workers, Fly.io |

vs Pull CDN vs Push CDN

| Feature | Pull CDN | Push CDN |
|--------------|---|---|
| How it works | CDN pulls content from the origin server on-demand | You upload content manually to the CDN |
| Storage | Origin is the source of truth | CDN stores the content permanently |
| Ideal for | Dynamic or frequently updated content | Large, rarely changing assets (e.g., videos) |
| Caching | Cached after first user request | Already cached because it was pre-uploaded |
| Example Flow | 1. User requests 2. CDN checks cache 3. If miss → fetch from origin | You upload content → CDN serves directly |
| Management | Easier, automatic | Manual effort to upload & update |

📦 Pull CDN (Most Common)

- You serve your site normally, and the CDN **caches** content the first time a user requests it.
- The next user nearby gets it from the **edge cache**.

✓ Pros:

- Simple setup

- Good for **blogs, websites, web apps**
- No need to manage content separately

Cons:

- First request might be slow (cache miss)
- Heavier load on origin server during cache expiry

Push CDN

- You **upload your content to the CDN** (via FTP, API, or dashboard), and it stays there.
- The CDN doesn't query your origin — it **serves only what you upload**.

Pros:

- Great for **large files** (videos, installers, PDFs)
- No hit to origin server
- Content always available at edge

Cons:

- Manual upload & versioning
- Not ideal for frequently updated content

Example Use Cases

| Use Case | Recommended CDN Type |
|------------------------------|----------------------|
| Blog or CMS site | Pull CDN |
| SaaS Web App (React/Next.js) | Pull CDN |
| Video Hosting / eBooks | Push CDN |
| Game file downloads / Assets | Push CDN |

TL;DR:

Pull CDN = "Just cache what users ask for."

Push CDN = "I'll upload everything ahead of time."

Bloom Filter

Purpose: Quickly check if an element *possibly exists* in a set (with false positives but no false negatives).

Use Case

- Membership testing (e.g., checking if a URL has been visited, a username already exists, or a password is in a breach list).
- Used in databases, caches, and search engines.

Core Idea

- A bit array of size m initialized to 0.
- k different hash functions.
- To insert an element: Hash it with k functions and set the bits at those indices to 1.
- To check membership: Hash again and check if all k positions are 1.
 - If *not*, element definitely not present.
 - If *yes*, it might be present (false positive possible).

Properties

- **False positives:** Yes
 - **False negatives:** No
 - **Space-efficient**
 - **Cannot remove elements** (unless you use a Counting Bloom Filter variant)
-

Count-Min Sketch

Purpose: Estimate the frequency of elements in a stream (with overestimations, but never underestimates).

Use Case

- Counting hashtags in real-time on Twitter.
- Counting page hits in a CDN.
- Network traffic analysis.

Core Idea

- A 2D array (table) with d rows and w columns (width), where each row has a hash function.
- To increment count of an item: hash it using each function and increment the respective positions.
- To get approximate frequency: take the **minimum** of all hashed values.

Properties

- **Overestimates** counts (due to collisions), but never underestimates.
- Memory efficient for approximate frequency tracking in data streams.
- Supports mergeability (useful in distributed systems).

Bloom Filter vs Count-Min Sketch

| Feature | Bloom Filter | Count-Min Sketch |
|------------------|---|---|
| Purpose | Membership test | Frequency estimation |
| False Positives | Yes | Yes (in frequency) |
| False Negatives | No | No |
| Insert Operation | Set bits | Increment counters |
| Query Type | Exists or not | Approx. frequency count |
| Space Efficiency | High | High |
| Removal Support |  (standard) /  (counting) |  |

40. Paxos – Consensus over Distributed Hosts

-  Leader Election

Paxos – Consensus in Distributed Systems

Paxos is a family of protocols for **reaching consensus** in a network of **unreliable processors (nodes)** — even when some fail or messages are lost/delayed. It's one of the **foundations of distributed systems** like databases, file systems, and cluster management.

Problem Paxos Solves:

| How do multiple nodes agree on a single value (e.g., leader, transaction log entry, config value) — even with failures?

Real-world Analogy:

Imagine a group of people trying to **agree on a meeting time**. Some may be slow to respond, others may not respond at all, and messages might be lost — yet they need a **reliable agreement**.

Goal:

- Ensure **only one value is chosen** (consistency).
 - Ensure **a value is eventually chosen** (liveness).
-

Paxos Roles:

1. Proposer

- Suggests values to be chosen.

2. Acceptor

- Votes on values. A **majority** must agree.

3. Learner

- Learns the final agreed value.
-

Basic Paxos Flow:

1. Prepare Phase

- Proposer sends a `prepare(n)` request with a unique proposal number `n` to acceptors.
- Acceptors reply with the highest proposal they've accepted (if any).

2. Promise

- If `n` is higher than anything seen before, the acceptor **promises not to accept proposals < n**.

3. Accept Phase

- Proposer sends `accept(n, value)` — using the value of the highest previous proposal, or a new one.

4. Accepted

- Acceptors accept it if they haven't promised a higher `n` already.

5. Consensus

- Once a **majority** of acceptors accept the same value, it's chosen. Learners are informed.
-

Properties:

| Property | Description |
|-----------------|--|
| Safety | Only one value can be chosen. |
| Liveness | Eventually a value will be chosen (with retries). |
| Fault Tolerance | Works even if some nodes crash or messages are lost. |

Paxos is used in:

- Google's **Chubby** (lock service)
- **etcd** (used by Kubernetes)
- Apache **ZooKeeper** (ZAB = modified Paxos)
- Raft (inspired by Paxos but more understandable)

Why is Paxos considered "hard to understand"?

Because:

- Lots of moving parts (roles, messages, majority quorum logic)
- Subtle edge cases
- Asynchronous assumptions (messages can be delayed/lost)

Problem Paxos Solves

In a **distributed system**, nodes may:

- Crash
- Restart
- Be slow or unresponsive
- Have **no shared memory or clock**

 How do all nodes **agree on a single value** (like a config, DB write, or leader election), even if some fail?

 **Answer:** Use Paxos to **achieve consensus**.

Core Idea

Paxos ensures that:

1. **Only one value** is chosen.

2. No two nodes disagree.
 3. Nodes may fail and recover without breaking consensus.
-

Roles in Paxos

| Role | Responsibility |
|----------|--|
| Proposer | Suggests a value to be agreed upon |
| Acceptor | Decides whether to accept a proposed value |
| Learner | Learns the final value once consensus is reached |

A node can take on multiple roles.

Basic Paxos Steps

1. Prepare Phase

- Proposer sends a **prepare request** with a unique number n to a majority of Acceptors.
- Acceptors respond with:
 - "OK" if $n >$ any previous prepare number
 - Their last accepted proposal, if any

2. Promise

- Acceptors promise **not to accept** proposals numbered less than n .

3. Propose Phase

- Proposer sends a **propose(n , v)** to the acceptors.
 - If an acceptor promised n , it accepts and sends back an acknowledgment.

4. Accept

- Once a **majority accepts**, the value is **chosen**.

5. Learn

- Learners are notified of the chosen value.
-

Challenges

- Multiple proposers can cause **conflicts**.

- Delays or retries are common.
 - **Implementation is hard** and verbose, though conceptually elegant.
-

Where It's Used

- Google's **Chubby lock service**
 - Microsoft's **Zookeeper alternatives**
 - Foundations of **Raft**, which is more understandable
-

Paxos vs Raft

| Aspect | Paxos | Raft |
|------------|----------------------|--|
| Complexity | Harder to understand | Easier, more practical |
| Roles | Flexible | Fixed leader role |
| Adoption | Mostly academic | Widely used in real-world (e.g., etcd, Consul) |

40. Virtual Machines & Containers

Virtual Machines (VMs)

◆ What is a VM?

A **VM** is a full emulation of a physical computer, running its own operating system (OS) on **virtualized hardware**.

A **hypervisor** (also known as a **Virtual Machine Monitor** or **VMM**) is software that allows multiple **virtual machines (VMs)** to run on a single **physical machine (host)** by **abstracting and managing hardware resources**.

Architecture:



✓ Pros:

- **Strong isolation:** Each VM has its own OS.
- Can run **different OSes** (e.g., Linux VM on Windows).
- Ideal for **legacy app support** and **full environment separation**.

✗ Cons:

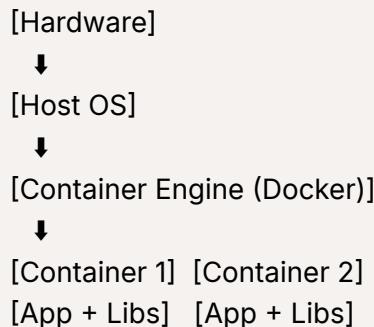
- **Heavyweight:** Each VM includes a full OS → more memory & CPU usage.
- **Slower startup time.**
- **Less efficient resource usage.**

📦 Containers

◆ What is a Container?

A **container** is a lightweight, portable package that bundles an app and its dependencies, **sharing the host OS kernel**.

⚙️ Architecture:



✓ Pros:

- **Lightweight & fast:** No separate OS → fast boot times.
- Better **resource efficiency**.
- Easier to **scale and deploy** (especially in microservices).
- Works well with **CI/CD** and orchestration tools (e.g., Kubernetes).

✗ Cons:

- **Weaker isolation** than VMs (kernel is shared).
- All containers run on the **same OS type**

- Slightly more **security-sensitive**.
-

VMs vs Containers

| Feature | Virtual Machines | Containers |
|----------------|------------------------|-----------------------------|
| OS Isolation | Yes (separate OS) | No (shared kernel) |
| Startup Time | Slow (minutes) | Fast (seconds) |
| Resource Usage | Heavy | Lightweight |
| Portability | Medium | High |
| Security | Stronger (isolated OS) | Slightly weaker (shared) |
| Use Case | Legacy, full-stack | Cloud-native, microservices |

Real-World Use Case Example:

- **VM**: Hosting multiple apps that require different OSes (e.g., Windows app + Linux backend).
- **Container**: Running microservices with quick deployments (e.g., in a Kubernetes cluster).

What is a Hypervisor?

A **hypervisor** (also known as a **Virtual Machine Monitor** or **VMM**) is software that allows multiple **virtual machines (VMs)** to run on a single **physical machine (host)** by **abstracting and managing hardware resources**.

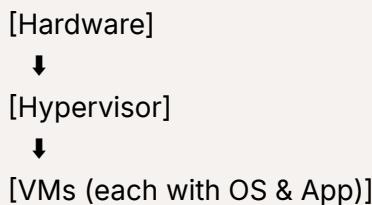
Types of Hypervisors

◆ Type 1 – Bare Metal Hypervisor

- Runs **directly on hardware** without a host OS.
- Used in **enterprise data centers** and cloud infrastructure.

Examples:

- VMware ESXi
- Microsoft Hyper-V
- Xen
- KVM (Linux Kernel-based VM)



✓ Pros:

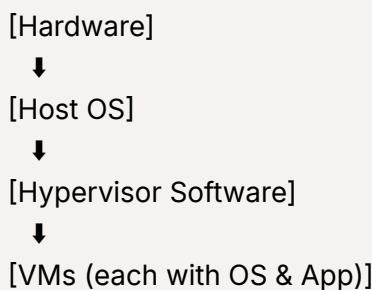
- Better performance
- More secure (less overhead)

◆ Type 2 – Hosted Hypervisor

- Runs **on top of a host operating system** like any other software.
- Used for **development, testing, and personal machines.**

Examples:

- VMware Workstation
- Oracle VirtualBox
- Parallels Desktop



✓ Pros:

- Easy to set up
- Great for desktop use

✗ Cons:

- Slightly slower (more layers)
- More vulnerable to host OS issues

🧩 What Does a Hypervisor Do?

- **Allocates CPU, memory, and disk** to each VM
 - **Manages VM isolation and security**
 - Handles **VM startup, pause, resume, shutdown**
 - Allows **live migration** (move VM from one host to another)
 - Supports **snapshots** (save/restore states)
-

Hypervisor vs Container?

| Feature | Hypervisor (VM) | Container (Docker) |
|-----------------|------------------------|--------------------------|
| OS per instance | Yes | No |
| Startup speed | Slow | Fast |
| Isolation | Strong (OS-level) | Moderate (process-level) |
| Resource usage | High | Low |
| Ideal for | Full OS virtualization | Microservices / scaling |

42. Publisher–Subscriber or Queue

Publisher–Subscriber (Pub/Sub) and **Queue-based messaging** are both messaging patterns used for **asynchronous communication** between services. Here's how they differ:

Publisher–Subscriber (Pub/Sub)

◆ What is it?

In Pub/Sub, **publishers** send messages to a **topic**, and **subscribers** receive messages from the topic. Multiple subscribers can receive the **same message**.

Flow:

```
[Publisher] → [Topic/Channel] → [Subscriber A]
                           → [Subscriber B]
```

Key Characteristics:

- **Fan-out** pattern
- Subscribers receive **copies** of the same message
- Loose coupling between producers and consumers
- Useful for **broadcasting events**

Use Cases:

- Email/SMS notifications
- Audit logs
- Event-driven architectures (e.g., Microservices)
- Real-time updates (e.g., chat, feed updates)

Tools:

- Google Cloud Pub/Sub
- AWS SNS
- Apache Kafka (Pub/Sub style)
- Redis Pub/Sub

Queue-based Messaging

What is it?

In message queues, **producers** send messages to a **queue**, and **one consumer** (at a time) processes each message.

Flow:

```
[Producer] → [Queue] → [Consumer A]  
                                [Consumer B] (but one msg goes to one consumer)
```

Key Characteristics:

- **Point-to-point** communication
- Each message is processed **only once**
- Guarantees order (in FIFO queues)
- Suitable for **task distribution**

Use Cases:

- Background jobs (e.g., image processing, email sending)
- Task queues in microservices
- Buffering high-throughput systems

Tools:

- AWS SQS
 - RabbitMQ
 - Celery (Python task queues)
 - Apache ActiveMQ
-

vs Comparison

| Feature | Pub/Sub | Queue |
|------------------|---------------------------|------------------------------|
| Message delivery | Multiple subscribers | One consumer per message |
| Use case | Broadcast / events | Task distribution |
| Coupling | Looser | Tighter |
| Example | Kafka, SNS, Redis Pub/Sub | RabbitMQ, SQS, Celery |
| Message loss | Possible if no subscriber | Usually stored till consumed |

Rule of Thumb:

- Use **Pub/Sub** when **multiple services need to know** about an event.
- Use **Queues** when **only one service needs to act** on a task.

43. **Fan-Out vs Fan-In** — two commonly used messaging and processing patterns in **distributed systems** and **asynchronous architectures**.

Fan-Out

Definition:

Fan-out is when **one producer** (or event) **triggers multiple consumers** or services to act in parallel.

Analogy:

Think of a celebrity tweet being received by millions of followers — all at once.

Flow:

[Event / Producer]





[Producer A] [Producer B] [Producer C]

✓ Use Cases:

- Notification systems (SMS, email, push notifications)
- Microservices reacting to a single event (e.g., order placed → inventory, billing, shipping services all act)
- Serverless function triggers (e.g., AWS Lambda triggered by SNS)

🔧 Tech Examples:

- AWS SNS → Multiple AWS Lambdas
- Kafka topic with multiple consumers
- Redis Pub/Sub
- Google Pub/Sub

⟳ Fan-In

◆ Definition:

Fan-in is when **multiple producers** send data or tasks into **one single consumer** or aggregator.

🧭 Analogy:

Like a restaurant with many waiters (producers) sending orders to one kitchen (consumer).

⟳ Flow:



[Producer A]
[Producer B] —→ [Single Consumer / Aggregator]
[Producer C]

✓ Use Cases:

- Logging and monitoring (many services push logs to one processor)
- Aggregating metrics or analytics from multiple services
- Multiple IoT devices sending data to a central hub

- Stream processing

Tech Examples:

- Apache Flink/Kafka Streams aggregating data
- AWS Kinesis stream consumers
- Logstash aggregating logs from multiple servers

Fan-Out vs Fan-In

| Feature | Fan-Out | Fan-In |
|-------------|---------------------------------------|--------------------------------------|
| Direction | 1 → many | Many → 1 |
| Use Case | Parallel event processing | Aggregation, collection |
| Examples | Notifications, multi-service triggers | Logging, analytics, batch processing |
| Scalability | Scales with consumers | Scales with producers |

Combined Patterns (Fan-Out + Fan-In):

In real-world architectures, you often use both:

- **Fan-out** to trigger parallel microservices.
- Each microservice **fans in** data to a common processing unit (like a data warehouse or reporting service).

43. 

MapReduce is a programming model and processing technique used to handle **large-scale data processing** across distributed systems. It was originally developed by Google to process massive datasets (think terabytes or petabytes) efficiently using a distributed cluster of machines.

Core Idea

The **MapReduce model** consists of two key functions:

- **Map:** Process and transform input data into key-value pairs.
- **Reduce:** Aggregate or combine values with the same key.

Workflow

Example Task: Count the frequency of words in a large text dataset

1. Map Step

- Input: Split raw data (text files) into chunks
- Task: Emit key-value pairs

```
Input: "cat bat cat dog"  
Output: ("cat", 1), ("bat", 1), ("cat", 1), ("dog", 1)
```

2. Shuffle & Sort Step

- Intermediate step (handled by the framework)
- Groups values by key:

```
Grouped: ("bat", [1]), ("cat", [1, 1]), ("dog", [1])
```

3. Reduce Step

- Aggregates all values per key:

```
("bat", 1), ("cat", 2), ("dog", 1)
```

Key Characteristics

| Feature | Description |
|-----------------------|---|
| Distributed | Splits and distributes work across multiple nodes |
| Fault-Tolerant | Automatically reruns failed tasks |
| Scalable | Efficiently processes data across 1000s of nodes |
| Batch-Oriented | Good for offline, high-throughput jobs, not low-latency |

Technologies Using MapReduce

- **Hadoop MapReduce** – The most famous open-source implementation
- **Google's internal MapReduce**
- **Spark** (uses more optimized DAG execution but inspired by MapReduce)
- **Amazon EMR, Azure HDInsight**

Real-World Use Cases

- Indexing the web (Google Search)
 - Log analysis
 - Data warehousing/reporting
 - Social media analytics
 - Recommendation engines
-

Summary

| Phase | Purpose |
|---------|---|
| Map | Split & convert data into key-value pairs |
| Shuffle | Organize intermediate data |
| Reduce | Aggregate values per key |

44. [Multithreading, Concurrency, Locks, Synchronization, CAS](#)

Multithreading

- **Definition:** Running multiple threads within a single process.
 - **Goal:** Achieve parallelism and better CPU utilization.
 - **Use Case:** Web servers (handling multiple requests), background tasks.
-

Concurrency vs Parallelism

| Concurrency | Parallelism |
|--|--|
| Multiple tasks start, run, and complete in overlapping time periods | Multiple tasks running simultaneously |
| Doesn't mean at the same time | Actually runs at the same time |
| Think: multitasking | Think: multiprocessing/multicore CPUs |

Locks

- **Purpose:** Prevent multiple threads from accessing **shared resources** at the same time.
 - **Types:**
 - **Mutex** (mutual exclusion): Only one thread can acquire the lock.
 - **Reentrant Lock**: Same thread can re-acquire the lock.
 - **Read-Write Lock**: Multiple reads allowed, only one write.
-

Synchronization

- A mechanism (like `synchronized` keyword in Java) to ensure only one thread accesses a resource at a time.
- **Prevents:** race conditions, inconsistent data.

Example:

```
synchronized void increment() {  
    count++;  
}
```

CAS (Compare-And-Swap)

- **Lock-free synchronization technique**
- Used heavily in **concurrent data structures** and **atomic operations**.

How CAS works:

1. Read a value from memory.
2. Check if it's the expected value.
3. If yes → update to new value.
4. If not → retry.

Java Example (using `AtomicInteger`):

```
AtomicInteger counter = new AtomicInteger(0);  
counter.compareAndSet(0, 1); // If value is 0, set to 1
```

When to Use What?

| Scenario | Use |
|-------------------------------|--|
| Simple mutual exclusion | <code>synchronized</code> or Mutex |
| High-concurrency structures | CAS or <code>Atomic*</code> classes |
| Multiple readers, few writers | ReadWriteLock |
| Avoiding blocking | Lock-free (CAS), non-blocking algorithms |

Summary

| Term | Purpose |
|-----------------|--|
| Multithreading | Use CPU better by running threads |
| Concurrency | Execute multiple tasks in overlapping time |
| Locks | Prevent data races via mutual exclusion |
| Synchronization | Coordination to access shared resources |
| CAS | Lock-free atomic updates |

45. Garbage Collection

Garbage Collection is the process of **automatically reclaiming memory** occupied by objects that are no longer in use, preventing memory leaks and improving application stability.

Why Garbage Collection?

- In languages like **Java, Python, Go, and C#**, memory is allocated on the heap for objects.
 - Once an object is no longer **reachable**, the GC reclaims that memory.
 - Manual memory management (like in C/C++) is error-prone and leads to issues like:
 - Memory leaks
 - Dangling pointers
 - Double frees
-

Basic GC Concepts

| Term | Meaning |
|-----------------------|---|
| Reachability | An object is "reachable" if it can be accessed from root variables (e.g., stack, globals) |
| Heap | Area of memory where objects are allocated |
| GC Root | Variables or references always accessible (e.g., stack frames, static fields) |
| Stop-the-world | GC halts application threads during collection |

How GC Works (in general)

1. **Mark:** Traverse all reachable objects starting from GC roots.
 2. **Sweep:** Remove all unreachable objects.
 3. **Compact (optional):** Move live objects to eliminate fragmentation.
-

GC Algorithms

1. Reference Counting

- Each object has a counter; incremented/decremented on assignment.
 - When count reaches zero, object is deleted.
 -  Issue: cannot handle **cyclic references**.
-

2. Tracing GC (Mark and Sweep)

- Used in Java, Go, etc.
 - Marks reachable objects → Sweeps the rest.
 - Variants:
 - **Stop-the-world** (simpler)
 - **Incremental**
 - **Concurrent GC** (minimal pause)
-

3. Generational GC (used in Java)

- **Young Generation:** Where new objects go (frequent collections).
 - **Old Generation:** Long-lived objects.
 - Assumption: most objects die young.
 - **Minor GC:** Collects young generation.
 - **Major GC / Full GC:** Collects old generation too.
-

Java GC Implementations

| GC Type | Best Use Case | Notes |
|------------------------------------|--------------------|-----------------------|
| Serial GC | Small applications | Single-threaded |
| Parallel GC | Multi-core systems | High throughput |
| CMS (Concurrent Mark Sweep) | Low pause time | Deprecated in Java 14 |
| G1 GC (Garbage First) | Balanced needs | Default since Java 9 |
| ZGC / Shenandoah | Very low latency | Java 11+ |

GC Tuning (Java Example)

```
-XX:+UseG1GC -Xms512m -Xmx2g -XX:MaxGCPauseMillis=100
```

In Python

- Uses **reference counting + cycle detector** (in `gc` module).
- Example:

```
import gc
gc.collect()
```

🔥 Benefits

- Automatic memory management
- Avoids common C/C++ issues

⚠️ Trade-offs

- **Pause times**
- **Less control** over memory
- Can impact **latency** or **throughput** if not tuned properly

46. `Idempotency`

🔄 Idempotency — Safe to Retry Without Side Effects

Idempotency means that **performing the same operation multiple times has the same effect as performing it once**. It's a key concept in system design, especially for APIs and distributed systems where retries are common.

📘 Real-World Analogy

- 🔒 Unlocking a door is **idempotent**: Doing it once or five times doesn't change the state — the door is still unlocked.
- 💡 Turning on a light switch *without checking* is **not idempotent** if each press toggles the light. If the light starts off, pressing once turns it on, pressing again turns it off — not idempotent.

✓ Examples of Idempotent HTTP Methods

| HTTP Method | Idempotent? | Example |
|---------------------|-------------|--|
| <code>GET</code> | ✓ Yes | Repeated calls fetch same resource |
| <code>PUT</code> | ✓ Yes | Updating a user profile with the same data repeatedly |
| <code>DELETE</code> | ✓ Yes | Deleting an already-deleted item returns the same result |
| <code>POST</code> | ✗ No | Each request typically creates a new resource |

To make POST idempotent, clients often send a unique idempotency key.

Why is Idempotency Important?

-  **Retries**: In case of network failures, clients may retry. Idempotent operations avoid duplication or corruption.
 -  **Distributed systems**: With async processing, eventual consistency, and retries, idempotency ensures consistency.
-

Implementing Idempotency

Example: Payment API

```
POST /charge
Headers:
Idempotency-Key: 12345
Body:
{ "amount": 100, "userId": 1 }
```

- Server checks if `Idempotency-Key` `12345` has already been used.
 - If yes, it **returns the same result** as before — no duplicate charge.
-

Where You'll See Idempotency Used

- Payment systems (Stripe, Razorpay, etc.)
 - Booking/reservations
 - Message queues
 - Distributed job schedulers
 - REST APIs & Microservices
-

Tips

- Store `idempotency_key → result` mapping with TTL (time-to-live).
- Ensure **same input + same key → same output**.
- Use **client-generated UUIDs or timestamps** as keys.

47. [Reverse Proxy](#)

Reverse Proxy — The Middleman That Works for Your Servers

A **reverse proxy** is a server that sits **in front of one or more backend servers**, intercepting client requests and forwarding them to the appropriate server. It then sends the server's response back to the client as if it originated from the proxy itself.

Reverse vs Forward Proxy

| Proxy Type | Sits In Front Of | Purpose | Example |
|----------------------|------------------|---|--|
| Forward Proxy | Clients | Access control, anonymity | Used at offices or by VPNs |
| Reverse Proxy | Servers | Load balancing, caching, TLS offloading | Used by companies like Google, Netflix, etc. |

What Can a Reverse Proxy Do?

1. Load Balancing

Distributes traffic across multiple servers for scalability and availability.

- e.g., [Nginx](#), [HAProxy](#), [Envoy](#)

2. SSL Termination

Handles HTTPS → HTTP conversion to offload encryption overhead.

3. Caching

Stores frequently accessed content closer to users for faster responses.

4. Compression

Reduces payload size (gzip, Brotli) to speed up delivery.

5. Security

Hides the backend structure from users, filters malicious requests.

6. Routing

Forwards requests to different microservices based on URL paths or headers.

7. Authentication / Authorization

Handles identity before passing requests to backend.

Popular Reverse Proxies

- **Nginx**
- **Apache HTTP Server**
- **HAProxy**
- **Envoy**

- **Traefik**
 - **Cloudflare** (acts as a global reverse proxy & CDN)
-

Diagram

```
Client —► Reverse Proxy —► Backend Server 1
          |           ↘► Backend Server 2
          |           ↙ handles caching, SSL, load balancing, etc.
```

Real-World Example

Say you're building a site like Amazon.

- Clients send requests to `amazon.com`
- Reverse proxy (e.g. Nginx) receives the request
 - If path is `/search`, it forwards to the **search service**
 - If path is `/checkout`, it forwards to the **order service**
- Handles all HTTPS, rate limiting, and load balancing logic

Forward Proxy — The Gatekeeper for Clients

A **forward proxy** is a server that sits **between a client and the internet**, acting on behalf of the client. It forwards client requests to the destination server **anonymously or with added control**, then returns the server's response back to the client.

Real-Life Analogy

Think of a forward proxy like a **valet**:

You tell the valet where to go (a website), and they go get it for you — hiding your identity from the destination.

How It Works

```
Client —► Forward Proxy —► Target Website (e.g., google.com)
          (hides client's IP, filters access)
```

Use Cases of Forward Proxy

| Use Case | Description |
|---|--|
|  Anonymity | Hides the client's IP address (used in VPNs or private browsing) |
|  Geo-Restriction Bypass | Access content not available in your country |
|  Content Filtering | Used in organizations to block certain websites (e.g., social media) |
|  Caching | Caches repeated requests to save bandwidth |
|  Authentication | Can require login credentials to allow internet access |
|  Monitoring/Logging | Companies use it to track employee browsing activity |

Example

Imagine a school that uses a **forward proxy**:

- All students' internet traffic first goes through the proxy.
- The proxy **blocks YouTube and Facebook**, allows only study-related sites.
- It **logs activity** and **caches** common requests to speed things up.

Tools / Technologies

- **Squid Proxy**
- **Apache Forward Proxy**
- **Proxy**
- **VPN services** (essentially forward proxies with encryption)

Quick Comparison: Forward vs Reverse Proxy

| Feature | Forward Proxy | Reverse Proxy |
|-----------------|-------------------------------|------------------------------|
| Who it protects | The client | The server(s) |
| Main use case | Anonymity, filtering, caching | Load balancing, SSL, routing |
| Request flow | Client → Proxy → Internet | Internet → Proxy → Server |

48. Network Procols: TCP, WebSocket, HTTP, etc.

1. TCP (Transmission Control Protocol)

Type: Transport Layer (Layer 4)

Connection: Connection-oriented

 Reliable, ordered, and error-checked delivery

 Ensures packets arrive **in order**

 Slightly slower due to overhead

 Use Cases: HTTP, HTTPS, FTP, SSH, SMTP

2. UDP (User Datagram Protocol)

Type: Transport Layer

Connection: **Connectionless**

 Faster, lightweight

 Unreliable, packets can be lost or out of order

 Use Cases: Video streaming, online games, VoIP, DNS

3. HTTP (HyperText Transfer Protocol)

Type: Application Layer (Layer 7)

Based on: TCP

Stateless & Request-Response model

- **HTTP/1.1:** One request per connection (keep-alive supported)
- **HTTP/2:** Multiplexing over a single connection (faster)
- **HTTP/3:** Based on QUIC (UDP), with faster handshakes

 Use Cases: Browsing websites, REST APIs

4. WebSocket

Type: Application Layer

Built on: TCP

Connection: Full-duplex (bidirectional), persistent

 Enables real-time communication

 Low latency

 Use Cases: Chat apps, live updates, gaming, trading dashboards

5. HTTPS (HTTP Secure)

HTTP + TLS/SSL encryption

 Secure communication

 Prevents eavesdropping & tampering

📦 Use Cases: Secure websites, banking, login pages

✉️ 6. FTP / SFTP / FTPS

- **FTP:** File Transfer Protocol (TCP)
- **SFTP:** SSH File Transfer Protocol (uses SSH)
- **FTPS:** FTP over SSL

📦 Use Cases: File transfers, backups

📞 7. SMTP / IMAP / POP3

Email protocols:

- **SMTP:** Sending emails
- **IMAP/POP3:** Receiving emails

🎮 When to Use What?

| Protocol | Real-Time | Reliability | Speed | Security | Use Case Example |
|-----------|-----------|-------------|-------|----------|------------------|
| TCP | ✗ | ✓ | 🚶 | ✗ | Web, FTP |
| UDP | ✓ | ✗ | 🏃 | ✗ | Games, VoIP |
| HTTP | ✗ | ✓ | 🚶 | ✗ | Webpages |
| HTTPS | ✗ | ✓ | 🚶 | ✓ | Banking |
| WebSocket | ✓ | ✓ | 🏃 | ✓ | Chat, Live Feed |

49. [Client-Server vs Peer-to-Peer Architecture](#)

🧭 Client-Server Architecture

📌 **Definition:** A central **server** provides services/data, and **clients** request them.

🔧 Characteristics:

- **Centralized control**
- Clients don't communicate with each other directly
- Server can be a performance bottleneck

✓ Pros:

- Easier to manage & secure
- Scalable (with load balancers)

- Centralized data storage

Cons:

- Single point of failure (if not redundant)
- Higher server cost

Examples:

- Web apps (Google, Facebook)
- Online banking
- REST APIs

Peer-to-Peer (P2P) Architecture

 **Definition:** All nodes (**peers**) are equal and can act as both client & server.

Characteristics:

- **Decentralized:** A **decentralized system** is a type of network where **control, data, and decision-making** are distributed across multiple nodes instead of relying on a single central authority.
- Each peer can share resources
- No central server needed

Pros:

- High fault tolerance
- Scales well with more peers
- Efficient for file distribution

Cons:

- Harder to secure/manage
- Data consistency is complex
- Latency can vary

Examples:

- Torrent clients (BitTorrent)
- Blockchain/cryptocurrency networks
- Skype (initial versions)

Summary Table:

| Feature | Client-Server | Peer-to-Peer |
|-------------------|-------------------------|----------------------|
| Architecture Type | Centralized | Decentralized |
| Communication | Client ↔ Server | Peer ↔ Peer |
| Scalability | Limited by server | Scales with peers |
| Fault Tolerance | Low (unless replicated) | High |
| Examples | Gmail, Amazon | BitTorrent, Ethereum |

50. Scale from 0 to Million Users

Phase 1: 0 – 1K Users (MVP / Validation)

- **Monolith First**
 - Single codebase / repo, one deployment unit
 - In-process calls (no network overhead)
- **Simple Stack**
 - One app server + one database instance
 - Use managed PaaS (Heroku, AWS Elastic Beanstalk)
- **Basic Caching**
 - In-app or local Redis/Memcached for hot data
- **Limited Monitoring**
 - Logs + basic metrics (CPU, memory, error rates)
- **CI/CD**
 - Automated build → deploy on push to main

Phase 2: 1K – 100K Users (Growth)

- **Scale the Web Tier**
 - Front it with a **Load Balancer** (L4 or L7)
 - Spin up multiple stateless app instances
- **Read Scalability**
 - Add **read replicas** for your database
 - Introduce **query caching** (Redis)

- **Edge Caching & CDN**
 - Serve static assets via CloudFront / Cloudflare
 - **Service Extraction**
 - Split out heavy modules (e.g., media processing) into separate services
 - **Enhanced Observability**
 - Prometheus + Grafana for metrics
 - Centralized logging (ELK / Loki)
 - **Automated Scaling**
 - Auto-scale app servers based on CPU or request latency
-

Phase 3: 100K – 1M Users (Scaling Up & Out)

- **Microservices / Service Mesh**
 - Break monolith into domain-aligned services
 - Use Istio/Linkerd for mTLS, traffic routing, circuit breaking
 - **Data Partitioning**
 - Shard your database (range or hash)
 - Introduce **CQRS** for complex read/write patterns
 - **Asynchronous Workflows**
 - Offload long tasks to queues (Kafka, RabbitMQ, SQS)
 - Use event-driven “fan-out” for notifications, analytics
 - **Global Footprint**
 - Deploy into multiple regions / AZs
 - Use global load balancers (Route 53 latency-based routing)
 - **Chaos Engineering**
 - Inject failures (Chaos Monkey) to harden reliability
-

Phase 4: 1M – 10M+ Users (Enterprise Grade)

- **Polyglot Persistence**
 - Mix relational, NoSQL, search (Elasticsearch), and graph stores as needed
- **Advanced Data Pipelines**

- Real-time streaming (Flink, Spark Streaming) for analytics
 - Data lakes and warehouses for BI (Snowflake, Redshift)
 - **Edge Compute & Personalization**
 - Run business logic at the edge (Lambda@Edge, Cloudflare Workers)
 - **Security & Compliance**
 - Zero Trust network, least-privilege IAM, continuous audit logging
 - GDPR/PCI/HIPAA controls where applicable
 - **Enterprise Observability**
 - Distributed tracing (Jaeger, Zipkin)
 - SLOs, error budgets, automated alerting & incident response
-

Core Practices Across All Phases

1. **Idempotent, Stateless Services**
 2. **Infrastructure as Code** (Terraform, Pulumi)
 3. **Immutable Deployments** (Docker, Kubernetes)
 4. **Exponential Backoff & Circuit Breakers** for resilience
 5. **Automated Backups & DR Drills**
 6. **Performance Budgeting** (95th-percentile latency targets)
-

TL;DR: Start simple, validate quickly, then layer in horizontal scaling, service decomposition, global distribution, and enterprise controls as demand rises.

51. Decentralized Systems

A **decentralized system** is a type of network where **control, data, and decision-making** are distributed across multiple nodes instead of relying on a single central authority.

Key Characteristics:

- **No single point of failure**
- **Autonomous nodes** that can make decisions independently

- **Distributed data & processing**
 - **Redundancy & fault-tolerance**
-

Advantages:

| Feature | Benefit |
|---|--|
|  Security | Harder to attack or take down the whole system |
|  Performance | Local processing reduces bottlenecks |
|  Availability | One node's failure doesn't halt the system |
|  Censorship Resistance | No central control makes censorship difficult |

Disadvantages:

| Feature | Challenge |
|--|---|
|  Complexity | Harder to design, test, and maintain |
|  Consistency | Maintaining data consistency is difficult |
|  Security Risks | Peer nodes may be malicious |

Examples of Decentralized Systems:

| Domain | Example |
|---------------|-------------------|
| File Sharing | BitTorrent, IPFS |
| Blockchain | Bitcoin, Ethereum |
| Communication | Matrix, Jami |
| Storage | Storj, Filecoin |

Decentralized vs Distributed vs Centralized:

| Feature | Centralized | Distributed | Decentralized |
|----------------|-----------------|-----------------------|----------------------|
| Control | Single point | Multiple, coordinated | Multiple, autonomous |
| Failure Impact | High | Medium | Low |
| Speed | Fast (at start) | Medium | Varies |

52. 

Apache **Kafka** is a **distributed event streaming platform** used to build real-time data pipelines and streaming applications. It's designed to be **fault-tolerant, scalable, and high-throughput**, making it ideal for handling massive streams of data.



Core Concepts

1. Producer

- Sends (publishes) data to Kafka topics.
- Data is usually sent in the form of key-value pairs.

2. Consumer

- Subscribes to topics and reads data.
- Can be part of a **consumer group**, allowing parallel data processing.

3. Topic

- Logical stream to which messages are sent.
- Each topic can have multiple **partitions** (enabling parallelism).

4. Partition

- Each topic is split into multiple partitions.
- A partition is an **ordered, immutable sequence of records**.
- Each message has an **offset**, a unique ID within a partition.

5. Broker

- A Kafka server that stores and serves data.
- A Kafka cluster consists of multiple brokers.

6. ZooKeeper / KRaft (newer)

- **ZooKeeper** is used for managing brokers, leader election, and metadata.
- Kafka is moving to **KRaft mode**, a built-in consensus mechanism to remove the need for ZooKeeper.

7. Retention

- Kafka stores messages for a configurable retention period (e.g., 7 days), even after they're consumed.



Use Cases

- **Real-time analytics** (e.g., monitoring, fraud detection)
- **Data integration** (e.g., stream data to Elasticsearch, S3)

- **Event sourcing**
 - **Log aggregation**
 - **Message queue** replacement
-

Kafka vs Traditional Message Queues

| Feature | Kafka | Traditional MQ (e.g., RabbitMQ) |
|-------------------|---------------------------------------|------------------------------------|
| Storage Model | Append-only log | In-memory/queue-based |
| Message Retention | Messages retained even after consumed | Messages deleted after consumption |
| Throughput | Very high | Lower than Kafka |
| Replayability | Yes (via offset) | No (once acknowledged) |
| Partitioning | Built-in | Not always supported |
| Durability | Disk-based, highly durable | Varies by implementation |

Key Kafka Design Choices

- **Immutable logs:** Messages are not updated, only appended.
 - **Decoupling:** Producers and consumers operate independently.
 - **Replay capability:** Consumers can rewind to any point via offset.
 - **Scalability:** Kafka scales horizontally by adding brokers and partitions.
 - **Fault tolerance:** Replication across brokers.
-

Delivery Semantics

- **At most once:** No duplicates, but possible data loss.
 - **At least once:** No data loss, but possible duplicates.
 - **Exactly once:** No duplicates and no data loss (needs idempotency and Kafka transactions).
-

Real-World Example

Imagine you run an e-commerce platform:

- **Producers:** Backend services publish events like "Order Placed," "Payment Received."
- **Kafka Topics:** Events go to topics like `orders`, `payments`.
- **Consumers:** Analytics engine consumes `orders` to update dashboards. Warehouse system consumes `orders` to fulfill them.

A **Message Queue** is a form of asynchronous service-to-service communication used in **distributed systems**, where messages are sent and stored in a queue until the **receiving service is ready** to process them.

Core Concepts

1. Producer (Sender)

- Sends messages to the queue.
- Doesn't need to wait for the consumer.

2. Queue (Buffer)

- Stores messages temporarily.
- Ensures **decoupling** between producers and consumers.

3. Consumer (Receiver)

- Pulls and processes messages from the queue.

Why Use Message Queues?

| Benefit | Description |
|---|--|
|  Decoupling | Producers and consumers don't need to run at the same time |
|  Retry & Resilience | If the consumer fails, the message can be retried |
|  Load Balancing | Multiple consumers can process messages in parallel |
|  Backpressure Handling | Queues act as a buffer for high traffic spikes |
|  Persistence | Messages can be stored until acknowledged |

Popular Message Queue Systems

| System | Type | Notes |
|---------------|--------------------|---|
| RabbitMQ | Traditional broker | Feature-rich, supports multiple protocols |
| Kafka | Distributed log | High-throughput, durable, replayable |
| Amazon SQS | Cloud-based | Fully managed, scalable |
| Redis Streams | In-memory | Good for lightweight pub-sub use |
| ActiveMQ | JMS (Java-based) | Older but still used in enterprise |

Delivery Semantics

- **At most once** – Message may be lost but never duplicated.

- **At least once** – Message will be delivered, possibly more than once.
 - **Exactly once** – Message is delivered once and only once (more complex to guarantee).
-

Example Use Case

 In an **e-commerce app**:

- Producer: Order Service sends `"order_created"` to the queue.
- Queue: Stores the event.
- Consumers:
 - Email service sends confirmation.
 - Inventory service adjusts stock.
 - Analytics service tracks purchase.

All run independently and reliably via the message queue.

Advanced Concepts

- **Dead Letter Queue (DLQ)**: Stores messages that couldn't be processed after several attempts.
- **Message Acknowledgement**: Ensures a message was processed successfully.
- **FIFO vs Pub/Sub**:
 - **Queue**: FIFO (First In, First Out)
 - **Pub/Sub**: Fan-out to multiple subscribers

54. 

A **Proxy Server** acts as an **intermediary** between a client (like a browser) and the internet. Instead of the client directly accessing the server, all requests go through the proxy.

What It Does:

1. **Receives a client request**.
 2. **Forwards it to the target server** (like a website).
 3. **Gets the response** from the target server.
 4. **Sends it back** to the client.
-

Benefits

| Feature | Description |
|-----------------------|--|
| Anonymity | Hides your IP address from the target site |
| Access Control | Restrict/block access to specific sites or content |
| Performance | Can cache responses to speed up repeated requests |
| Security | Acts as a firewall, filters out malicious content |
| Logging | Tracks and logs internet usage |

Types of Proxy Servers

| Type | Direction | Use Case Example |
|----------------------|-------------------|--|
| Forward Proxy | Client → Internet | Used in schools/offices to restrict or monitor browsing |
| Reverse Proxy | Internet → Server | Used by companies (e.g., Nginx) to protect backend servers |

Real-World Examples

- **Forward Proxy:** A company's network blocks YouTube; requests go to the proxy, which denies access.
- **Reverse Proxy:** A user hits `example.com`, but the request is routed to one of many backend servers by Nginx.

Advanced Uses

- **Load balancing** (e.g., with reverse proxies)
- **SSL Termination**
- **Rate limiting**
- **Geo-blocking**

55. `Storage Types:`

- `Block Storage`

Block Storage – Simple Explanation

Block storage is a type of data storage where data is stored in fixed-sized blocks. It's commonly used in cloud services, virtual machines, and databases.

Key Idea

Block storage **splits data into blocks**, each with a unique identifier, and stores them separately. These blocks can be **distributed** across multiple storage systems or devices and reassembled when needed.

How It Works

- A file (say a 1 GB video) is broken into smaller blocks (e.g., 4KB each).
 - Each block is stored individually and can be accessed or modified independently.
 - Your system (OS/VM) reassembles blocks when the file is read.
-

Characteristics

| Feature | Description |
|---|---|
|  Raw Storage | Doesn't care about file structure—just blocks |
|  High Performance | Optimized for read/write-intensive tasks like databases |
|  Persistent | Blocks stay even after reboots, unlike memory |
|  Custom File System | OS/VM formats and manages it as needed (e.g., ext4, NTFS) |

Block Storage vs Other Types

| Storage Type | Use Case | Access Type | Example Services |
|-----------------------|-------------------------------|------------------------|-------------------------------|
| Block Storage | Databases, VMs, boot volumes | Low-level blocks | AWS EBS, GCP Persistent Disks |
| File Storage | Shared folders, file systems | Files/Folders | NFS, SMB, Dropbox |
| Object Storage | Static assets, backups, media | Object (with metadata) | AWS S3, GCP Cloud Storage |

Use Cases

- **Databases** (PostgreSQL, MySQL)
- **Virtual machines** (OS disk)
- **Email servers**
- **Transactional systems**
- **File Storage**
- **Object Storage (e.g., S3)**
- **RAID**

RAID is a data storage virtualization technology that **combines multiple physical disks** into one logical unit to improve **performance, fault tolerance**, or both.

Why Use RAID?

- **Speed**: Combine disks for faster read/write
 - **Redundancy**: Prevent data loss in case of disk failure
 - **Scalability**: Expand storage seamlessly
-

Common RAID Levels

| RAID | Disks Required | Description | Pros | Cons |
|------|----------------|---|-------------------------------|---|
| 0 | ≥ 2 | Striping (split data across disks) | Fast read/write | <input checked="" type="checkbox"/> No fault tolerance (1 fail = data loss) |
| 1 | ≥ 2 | Mirroring (exact copies on disks) | Redundancy, fast read | <input checked="" type="checkbox"/> Expensive (50% storage used for copies) |
| 5 | ≥ 3 | Striping + Parity (recovery info) | Balanced speed + redundancy | <input checked="" type="checkbox"/> Slower writes (due to parity calc) |
| 6 | ≥ 4 | Like RAID 5, but 2 parities | Survives 2 disk failures | <input checked="" type="checkbox"/> Even slower writes than RAID 5 |
| 10 | ≥ 4 | RAID 1 + RAID 0 (mirror + stripe) | High performance + redundancy | <input checked="" type="checkbox"/> Costly (only 50% usable space) |

Key Concepts

- **Striping**: Divides data across disks → parallel reads/writes = speed boost
 - **Mirroring**: Duplicates data → redundancy = safety
 - **Parity**: Stores calculated data to rebuild lost blocks if a disk fails
-

Use Cases

-  RAID 0: Temporary fast storage (video editing)
-  RAID 1: Critical data needing backup (bank logs)
-  RAID 5/6: General-purpose servers, NAS
-  RAID 10: High-performance databases, enterprise systems

56. **File System:**

-  Google File System
-  HDFS

A **File System (FS)** is a method used by operating systems to **organize, store, retrieve, and manage data on storage devices** (like SSDs, HDDs, USB drives).

What Does a File System Do?

-  **Organizes** files into directories/folders
-  **Tracks** where each file is stored on disk
-  **Manages permissions** for security and access
-  **Maintains metadata** (file size, timestamps, etc.)

Common Types of File Systems

| File System | Used In | Features |
|-------------|-----------------------------|--|
| FAT32 | USBs, older Windows systems | Lightweight, cross-platform, limited to 4GB/file |
| NTFS | Windows | Supports large files, permissions, journaling |
| ext3/ext4 | Linux | Journaling, performance, large file support |
| APFS | macOS | Encryption, snapshots, crash protection |
| HFS+ | Older macOS | Predecessor to APFS |
| XFS, Btrfs | Linux servers | High performance, scalability, snapshots |

Key File System Concepts

- **Inodes:** Store metadata about files (size, permissions, timestamps)
- **Blocks:** Fixed-size chunks of data on disk; files are stored in them
- **Journaling:** Logs changes before they're made → helps recover from crashes
- **Mounting:** Linking a file system to a specific directory in the OS

Why It Matters in System Design?

- Affects **performance** (e.g., large file access speed)
- Influences **reliability** (journaling vs non-journaling)
- Impacts **scalability** and **concurrency** in distributed systems

57. 

Bloom Filter – Explained Simply

A **Bloom Filter** is a **space-efficient probabilistic data structure** used to test whether an element **is possibly in** a set or **definitely not in** the set.

Core Properties

- **Fast** lookups

- **Memory-efficient**
 - **No false negatives**
 - **Possible false positives**
-

🧠 How It Works

1. **Initialization:** You have a bit array of size m , all set to 0.
 2. **Hashing:** Use k independent hash functions.
 3. **Insertion:** For each element:
 - Hash it k times
 - Set the bits at the resulting k positions to 1
 4. **Query:** To check if an element exists:
 - Hash it k times
 - Check if all k bits are 1
 - If **yes** → "maybe in set"
 - If **no** → "definitely not in set"
-

📦 Example

```
Bit array: 0 0 1 1 0 1 0 0 1 0  
Element: "apple"  
Hash1("apple") → index 2  
Hash2("apple") → index 3  
Hash3("apple") → index 5  
→ All bits = 1 → "apple" might be in the set
```

🔥 Real-World Use Cases

- **Web caching:** Check if a URL is already cached
- **Database queries:** Avoid expensive lookups for missing keys
- **Distributed systems:** e.g., HBase, Cassandra, Bigtable
- **Email spam filters:** Flag known spam indicators

Let's say we have:

- A **bit array** of size 10:

0 0 1 1 0 1 0 0 1 0

- 3 **hash functions**:

- Hash1("apple") → 2
- Hash2("apple") → 3
- Hash3("apple") → 5

✓ Inserting "apple" into the Bloom Filter:

These hash indices (2, 3, 5) are set to 1.

Now, when we check if "apple" exists:

- Check bit at index 2 → 1
- Check bit at index 3 → 1
- Check bit at index 5 → 1

Since **all are 1**, the filter returns:

| ✓ "apple might be in the set"

! What if "banana" hashes to the same indices?

Even though we **never added "banana"**, the bits it hashes to may coincidentally already be set to 1 (by "apple" or others), so:

| ! "banana might be in the set" (→ false positive)

✗ If even one bit is 0, like:

- Index 2 → 1
- Index 3 → 1
- Index 7 → 0 ← not all 1

Then:

| ✗ Definitely not in the set

⚠️ Trade-Offs

| Pros | Cons |
|-----------------------------|--|
| Extremely memory efficient | Can return false positives |
| Very fast operations | Cannot delete items (in classic form) |
| Scales well with large data | Needs tuning (size of array & #hashes) |

58. Merkle Tree, Gossiping Protocol

A **Merkle Tree** (or hash tree) is a **binary tree** where each leaf node represents a hash of data, and non-leaf nodes represent hashes of their respective child nodes. It is primarily used to verify the integrity of large datasets in a decentralized manner, ensuring that no data has been tampered with. Merkle Trees are essential in **blockchain**, **distributed systems**, and **cryptography**.

How it works:

- Leaf Nodes:** Each leaf node holds a hash of a piece of data (e.g., a transaction or block).
- Non-leaf Nodes:** Each non-leaf node holds a hash of the concatenation of its two child nodes' hashes.
- Root Node:** The final hash at the top (the root) represents the hash of all the data below it. If any piece of data changes, the root hash will change, indicating data integrity has been compromised.

Advantages:

- Efficient Data Verification:** Only a small subset of the tree is needed to verify if a particular piece of data is part of the dataset.
- Security:** Even one bit change in the data will change the root hash.
- Scalable:** Useful in scenarios like blockchain, where verification of data consistency and integrity is critical.

Example:

If you have 4 transactions and you want to prove that a particular transaction exists in the set without revealing the whole dataset, you only need a few hashes from the tree instead of the entire dataset.

Gossiping Protocol:

The **Gossip Protocol** is a communication method used in distributed systems for **sharing state information** across the system. It is an efficient, probabilistic way to propagate information in a large, distributed network, ensuring that data eventually reaches all nodes, even in the presence of failures or partitions.

How it works:

- **Peers** periodically "gossip" or share their state information with a random peer.
- This process repeats until all nodes have received the information.
- Each peer does not need to know the entire network, just a few peers to gossip with, which helps scale the system.

Key Characteristics:

- **Eventual Consistency:** Gossip protocols provide eventual consistency rather than immediate consistency. Over time, all nodes will converge to the same state.
- **Fault Tolerance:** Even if some nodes are temporarily down or unreachable, the protocol ensures data will propagate when they come back online.
- **Scalable:** Because nodes communicate with random peers rather than all nodes, gossip protocols scale well in large systems.

Example:

- **Amazon DynamoDB** and **Cassandra** use gossip protocols to share information about node state (like availability and partitioning).
- **Peer-to-Peer Systems:** Gossip protocols help distribute and sync information across nodes in a decentralized manner.

Advantages:

- **Simple and Robust:** Gossip protocols are simple to implement and highly fault-tolerant.
- **Low Overhead:** Since each node only communicates with a small subset of peers, it reduces overhead and doesn't require central coordination.
- **Scalable:** Perfect for distributed systems and large-scale applications.

Use Cases:

- **Merkle Trees:** Blockchain (e.g., Bitcoin), distributed file systems (e.g., IPFS), and data verification in large datasets.
- **Gossiping Protocol:** Distributed databases (e.g., Cassandra, DynamoDB), cloud services, and peer-to-peer networks.

59. Caching:

- Cache Invalidation

This refers to **removing stale or outdated data** from the cache when the underlying source of truth changes.

Common Strategies:

- **Write-through**: Data is written to the cache and the database simultaneously.
- **Write-around**: Data is written only to the database, and the cache is updated only on read.
- **Write-back (Write-behind)**: Data is written to the cache first, and then written to the database asynchronously.
- **Time-based Expiry (TTL)**: Each item is assigned a time-to-live after which it's invalidated.

Goal: Ensure consistency between cache and source of truth (like DB).

- **Cache Eviction** (LRU, LFU, MRU)

Eviction occurs when the cache is full and old data needs to be removed to make space for new data.

Eviction Policies:

- **LRU (Least Recently Used)**

Evicts the item that was least recently accessed.

● Best for most scenarios with temporal locality.

- **LFU (Least Frequently Used)**

Evicts items that are used the least number of times.

⌚ Useful when some items are consistently accessed more.

- **MRU (Most Recently Used)**

Evicts the most recently used item first.

❗ Rarely used, but good in specific scenarios (e.g., user might not reuse the most recent item).

30. **Back of the envelope estimation**

Back-of-the-envelope estimation is a **quick, rough calculation** to estimate feasibility, scale, or system requirements without going deep into exact data. It's super handy in **system design interviews, capacity planning, or architecture discussions**.

✓ When Do You Use It?

- Estimate **storage needs**
- Estimate **requests per second (RPS)**
- Estimate **bandwidth or throughput**

- Estimate **latency impact**
 - Estimate **hardware cost or scaling factor**
-

📦 Example 1: Storage Estimation

"Design a photo-sharing app like Instagram."

Assume:

- Each user uploads 5 photos/day
- Each photo = 3 MB
- 1 million users

📌 Daily Storage Need:

$$5 \text{ photos / user} * 3 \text{ MB} * 1\text{M users} = 15,000,000 \text{ MB/day} = \sim 15 \text{ TB/day}$$

📌 Monthly:

$\sim 450 \text{ TB}$

⚡ Example 2: RPS Estimation

"Handle login requests at peak traffic."

Assume:

- 10M daily active users
- 20% log in during peak hour
- That's 2M logins in 1 hour

📌 Peak RPS:

$$2\text{M logins / 3600s} \approx 555 \text{ RPS}$$

🌐 Example 3: Bandwidth Estimation

"Each request returns a JSON response of ~100 KB. Peak RPS = 500.

📌 Bandwidth =

$$500 \text{ req/s} * 100 \text{ KB} = 50,000 \text{ KB/s} = \sim 50 \text{ MB/s}$$

$$= \sim 400 \text{ Mbps}$$

12 34 Typical Unit Conversions

| Unit | Value |
|------|-------|
| | |

| | |
|-----------|-------------------------|
| 1 KB | 1,000 bytes |
| 1 MB | 1,000 KB (or 1e6 bytes) |
| 1 GB | 1,000 MB (or 1e9 bytes) |
| 1 TB | 1,000 GB |
| 1 Gbps | ~125 MB/s |
| 1 million | 10^6 |

31. [Sharding \(Horizontal and Vertical\)](#)

Sharding: Horizontal vs Vertical

Sharding is a **database partitioning technique** used to scale large systems by distributing data across multiple machines.

Horizontal Sharding (aka Data Sharding)

Think: **Rows-based split** across different databases.

Example:

A `Users` table with millions of rows can be split by user ID:

| Shard 1 | Shard 2 | Shard 3 |
|-----------|------------|------------|
| User 1–1M | User 1M–2M | User 2M–3M |

Each shard holds **the same schema**, but **different data rows**.

Pros:

- Scales well with user growth
- Distributes load evenly
- Easy to add more shards

Cons:

- Complex joins across shards
- Cross-shard transactions are hard
- Requires a shard key strategy

Vertical Sharding (aka Functional Sharding)

Think: **Columns/services-based split** by feature/module.

Example:

Instead of one monolithic DB, you split by domain:

- **Auth DB** : handles login info
- **Profile DB** : handles user profiles
- **Orders DB** : stores purchases

Pros:

- Encapsulates functionality
- Teams can own independent modules
- Optimizes schema and performance per domain

Cons:

- Still limited by vertical scaling
- Not effective if one table is still huge
- Can lead to tight coupling between services

TL;DR

| Feature | Horizontal Sharding | Vertical Sharding |
|---------------------|-----------------------------|----------------------------|
| Split by | Rows / Data | Columns / Features |
| Use case | High volume of similar data | Different data domains |
| Example | Split Users by ID | Split DB into Auth, Orders |
| Joins across shards | Difficult | Less of an issue |
| Scaling | Better for big data | Better for complexity |

32. Partitioning

Partitioning is the process of **dividing a large database table** into smaller, manageable pieces (partitions), while keeping it logically as a single table to the application.

Why Partition?

- Improve **performance** (faster reads/writes)
- Enhance **maintainability** and **scalability**
- Reduce **index size**
- Easier **archiving** and **backups**

Types of Partitioning:

1. Horizontal Partitioning (Sharding)

- Split data **row-wise**
- Each partition contains a **subset of rows**

📌 Example:

```
Users_Part1 → User ID 1-1M  
Users_Part2 → User ID 1M-2M
```

● Common in distributed systems (scaling out).

2. Vertical Partitioning

- Split data **column-wise**
- Each partition contains a **subset of columns**

📌 Example:

```
UserCredentials(user_id, email, password)  
UserProfile(user_id, name, age, location)
```

● Reduces I/O if queries only need a few columns.

3. Range Partitioning

- Based on a **range of values**

📌 Example:

```
Orders_2022  
Orders_2023  
Orders_2024
```

● Great for time-series data or archiving.

4. List Partitioning

- Based on **discrete values**

📌 Example:

```
India_Customers  
USA_Customers  
UK_Customers
```

5. Hash Partitioning

- Based on a **hash function** (e.g., `hash(user_id) % N`)

✖ Randomly distributes data to avoid hot partitions.

● Good when data lacks a natural range/list.

vs Partitioning vs Sharding

| Feature | Partitioning | Sharding |
|------------|------------------------------|------------------------------------|
| Scope | Usually within one DB server | Spreads across multiple DB servers |
| Managed by | DB engine (e.g., PostgreSQL) | Application or custom logic |
| Use case | Moderate scaling | Massive scaling (millions+) |

53. `Replication, Mirroring`

Replication is the process of copying and maintaining **database objects** (like tables) across multiple **database servers**.

◆ Types of Replication:

1. Master-Slave Replication

- **Writes → Master, Reads → Slave(s)**
- Slaves replicate data asynchronously (or semi-synchronously) from master
- Used for read scaling and backup

2. Master-Master Replication

- Both nodes **can read and write**
- Conflicts must be handled (e.g., by conflict resolution strategies)

3. Synchronous vs Asynchronous

- **Synchronous**: Writes are confirmed only after replication
- **Asynchronous**: Writes are confirmed immediately; replication happens later

✓ Use Cases:

- Read-heavy applications
- Failover readiness
- Geographical distribution

🔍 Mirroring

Mirroring is typically a **1-to-1** copy of a **whole database** or disk block at the **storage level**, mostly used in traditional RDBMS and storage systems.

◆ Characteristics:

- **Real-time copy** of the database to another server
- The mirror server is in **standby** (no reads/writes)
- Fast failover: if primary fails, mirror takes over

❖ Database Example:

- SQL Server Database Mirroring
- RAID 1 (disk-level mirroring)

✓ Use Cases:

- High availability
- Disaster recovery
- Fast failover without data loss

vs Key Differences

| Feature | Replication | Mirroring |
|--------------------|--------------------------------|---------------------------------|
| Copy Level | Logical (rows, tables) | Physical (entire DB or storage) |
| Read/Write Support | Slaves often support reads | Mirror is usually read-only |
| Nodes | One or many slaves | Usually one mirror |
| Failover Support | Manual or semi-automated | Often automatic failover |
| Use Case | Scalability, read distribution | High availability, DR |

34. [Leader Election](#)

Leader election is the process of **choosing one node** (from among a group of distributed nodes) to act as a **coordinator or master** for a particular task. This ensures **consistency, coordination, and conflict resolution** across nodes.

🔧 Why is Leader Election Needed?

- In distributed systems (e.g., Kafka, Zookeeper, Kubernetes), multiple nodes work together.
- Only **one leader** should:
 - Accept writes
 - Manage consensus

- Coordinate worker nodes
 - Prevents conflicts and ensures fault tolerance.
-

Common Leader Election Algorithms

1. Bully Algorithm

- Each node has a unique ID.
- Node with the **highest ID becomes the leader**.
- If a node crashes, the next highest ID takes over.
- Simple, but assumes reliable communication.

2. Raft (Used in etcd, Consul)

- Nodes start in **follower** state.
- A follower that times out becomes a **candidate** and starts an election.
- Majority votes → becomes **leader**.
- Heartbeats maintain leadership.

3. Paxos

- More complex than Raft.
- Focuses on achieving consensus for each value (not just electing a leader).
- Used in systems like Google's Chubby.

4. Zookeeper

- Clients create **ephemeral sequential znodes**.
 - Node with the **lowest znode ID becomes the leader**.
 - If leader node dies (its znode vanishes), election re-runs.
-

Use Cases

| System | Leader Election Role |
|--------------------|--|
| Kafka | Leader broker for each partition |
| Zookeeper | Elects a master for coordination |
| Kubernetes | Leader controller-manager for API operations |
| Redis Sentinel | Elects a master instance |
| Raft-based systems | Leader handles client requests |

Key Properties

- **Fault-tolerant:** Can re-elect on failure
- **Consistent:** One leader at any time
- **Fast Recovery:** Minimizes downtime

35. [Indexing etc.](#)

Indexing is a technique used to **speed up data retrieval operations** on a database table at the cost of additional **writes and storage space**. It's analogous to a book's index — instead of scanning every page, you go directly to the relevant page.

Why Use Indexes?

- **Speeds up SELECT queries**
- Helps with:
 - WHERE clauses
 - JOINs
 - ORDER BY
 - GROUP BY
- **Trade-off:** Slower writes (INSERT, UPDATE, DELETE)

Types of Indexes

| Type | Description |
|------------------------|--|
| B-Tree Index | Default in most RDBMS; maintains sorted data. Ideal for range queries. |
| Hash Index | Faster for exact matches. Not good for range queries. |
| Bitmap Index | Efficient for low-cardinality columns (e.g., gender, country). |
| Composite Index | Index on multiple columns (e.g., <code>(last_name, first_name)</code>). |
| Full-text Index | Used for searching large text blocks. |
| Spatial Index | Optimized for spatial/geometric data (e.g., maps). |
| Covering Index | Index includes all data needed by the query (no need to access table). |
| Unique Index | Ensures all values in indexed column(s) are distinct. |

Indexing Concepts

| Concept | Meaning |
|-------------------|---------------------------------------|
| Index Scan | Using index to find data efficiently. |

| | |
|-----------------------------------|--|
| Sequential Scan | Scanning table row-by-row. Slower, but needed if index isn't useful. |
| Index Selectivity | How unique the values in an indexed column are. Higher = better index. |
| Clustered vs Non-clustered | Clustered = index defines row order on disk. Non-clustered = separate. |

⚠ When Not to Use Indexes

- Very frequent INSERTs/UPDATEs
- Small tables (sequential scan is faster)
- Low-cardinality columns (unless using Bitmap Index)

🔧 Additional "Etc." Related to Indexing

- **Vacuuming (Postgres)**: Clears dead tuples caused by updates/deletes.
- **Reindexing**: Rebuilds corrupted or bloated indexes.
- **Partial Indexes**: Index only rows that satisfy a condition (`WHERE is_active = true`)
- **Functional Indexes**: Index on expressions (e.g., `LOWER(name)`)

36. NGINX

🌐 NGINX – High Performance Web Server, Reverse Proxy & More

NGINX (pronounced "engine-x") is a **lightweight, high-performance web server** often used as a **reverse proxy, load balancer, HTTP cache, and media streaming server**. It's especially popular for handling **high concurrency**.

🚀 Common Use Cases

| Use Case | Description |
|----------------------------|--|
| Web Server | Serves static files (HTML, CSS, JS, images, etc.) |
| Reverse Proxy | Fowards client requests to backend servers (e.g., Node, Python, Java) |
| Load Balancer | Distributes requests to multiple backend servers using round-robin, etc. |
| SSL/TLS Termination | Handles encryption at edge and forwards plaintext to internal services |
| Caching | Caches responses to reduce load on application/backend |
| Rate Limiting | Prevents abuse or DoS attacks by limiting requests per user/IP |

🔄 Reverse Proxy Setup

```
nginx
CopyEdit
```

```

server {
    listen 80;

    location / {
        proxy_pass http://localhost:3000;
        proxy_set_header Host $host;
        proxy_set_header X-Real-IP $remote_addr;
    }
}

```

 This forwards all HTTP requests on port 80 to a local backend server running on port 3000.

Key Concepts

| Feature | Description |
|-------------------------------|--|
| Event-driven Model | Uses non-blocking I/O to handle thousands of simultaneous connections |
| Modules | Highly modular (Gzip, SSL, geoip, Lua scripting, etc.) |
| Worker Processes | Master process spawns worker processes to handle requests concurrently |
| Static File Efficiency | Extremely fast at serving static files directly from disk |

Load Balancing Algorithms in NGINX

```

upstream backend {
    server backend1.example.com;
    server backend2.example.com;
}

```

 You can customize balancing with:

- `round-robin` (default)
- `least_conn`
- `ip_hash`
- `weight` (weighted round robin)

SSL Termination Example

```

server {
    listen 443 ssl;
    ssl_certificate /etc/ssl/certs/server.crt;
    ssl_certificate_key /etc/ssl/private/server.key;

    location / {
        proxy_pass http://localhost:3000;
    }
}

```

Bonus: Static Caching Example

```

location /static/ {
    root /var/www/html;
    expires 7d;
    add_header Cache-Control "public";
}

```

37. Load Balancer

A **Load Balancer** is a component that **distributes incoming network traffic** across multiple servers (called **backend servers** or **nodes**) to ensure no single server becomes a bottleneck.

Why Use a Load Balancer?

| Benefit | Explanation |
|------------------------------|--|
| Scalability | Distributes load across many servers |
| High Availability | If one server goes down, traffic is rerouted to healthy ones |
| Fault Tolerance | Detects failures and removes them from the pool |
| Improved Performance | Reduces response time and prevents overloading |
| Zero Downtime Deploys | Traffic can be gradually shifted to new servers during deployments |

Load Balancer Types

| Type | Layer | Example Use |
|-----------|--------------------------------|--|
| L4 | Transport Layer (TCP/UDP) | Fast, works at IP level, good for any protocol |
| L7 | Application Layer (HTTP/HTTPS) | Smart routing based on URL, headers, cookies, etc. |

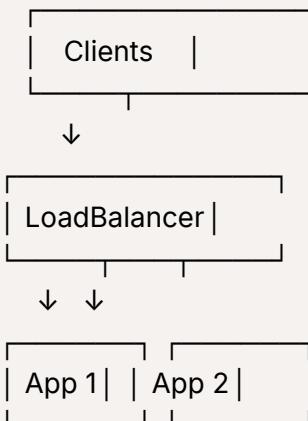
Common Load Balancing Algorithms

| Algorithm | Description |
|---------------------------|---|
| Round Robin | Evenly rotates requests among servers |
| Weighted RR | Some servers get more traffic (based on capacity) |
| Least Connections | Sends new request to server with fewest current connections |
| IP Hash | Client IP determines which server gets the request |
| Random | Requests are sent randomly to any server |
| Consistent Hashing | Useful for sticky sessions or cache locality |

Load Balancer in Action (Example with NGINX):

```
upstream myapp {  
    server app1.example.com;  
    server app2.example.com;  
}  
  
server {  
    location / {  
        proxy_pass http://myapp;  
    }  
}
```

Load Balancer Architecture



Examples of Load Balancers

| Provider | Type |
|---------------------------|--------------------|
| NGINX | Software (L4 & L7) |
| HAProxy | Software (L4 & L7) |
| AWS ELB/ALB/NLB | Cloud-based |
| Cloudflare Load Balancing | L7, global reach |

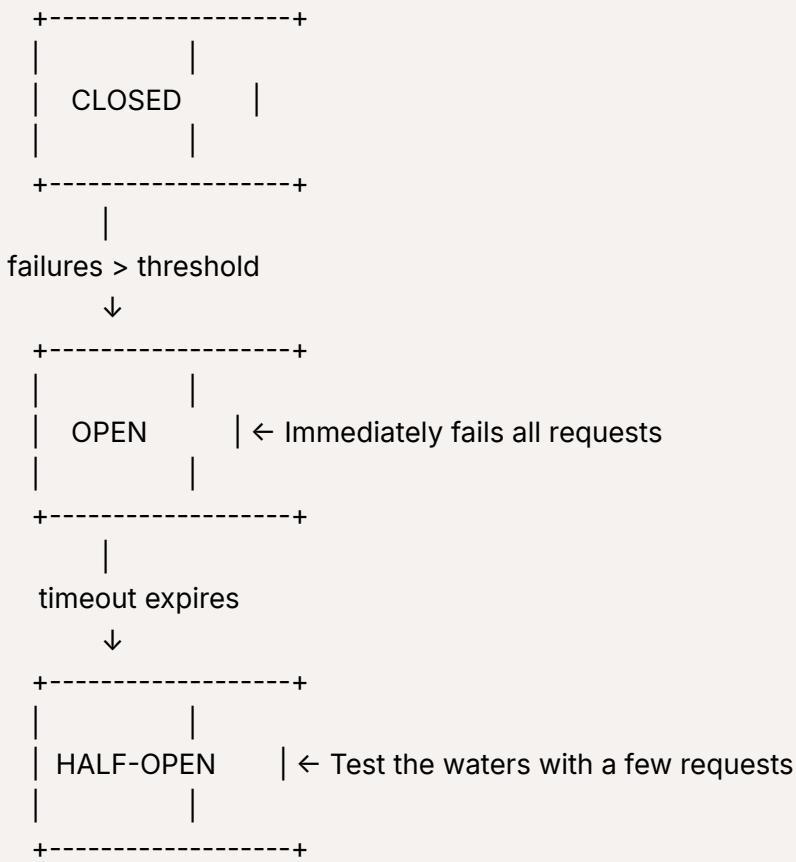
37. Circuit Breaker

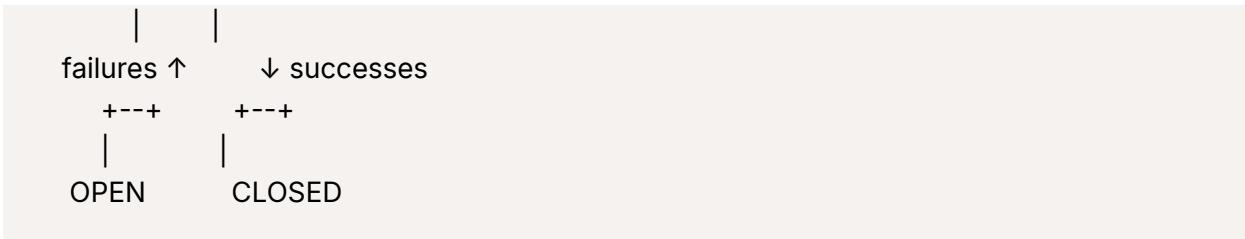
A **Circuit Breaker** is a **resilience design pattern** used in distributed systems to **detect failures** and **prevent continuous retries** on a failing service, which could lead to cascading failures and resource exhaustion.

Real-world Analogy

Like an **electrical circuit breaker** that cuts off power when there's a surge, a **software circuit breaker** stops calls to a failing service to give it time to recover.

States of a Circuit Breaker





✓ Behavior Summary

| State | Behavior |
|------------------|--|
| Closed | All requests pass through. Failures are monitored. |
| Open | All requests are failed fast (fallback logic is triggered). |
| Half-Open | A few test requests are allowed. If successful → go to Closed. If fail → back to Open. |

💡 Benefits

- Prevents **resource exhaustion** (threads, memory).
- Enables **fast failure** and **fallback logic** (e.g., return cached data).
- Allows the system to **self-heal** over time.

✍ Example (Pseudo-code in Java using Resilience4j)

```

CircuitBreaker circuitBreaker = CircuitBreaker.ofDefaults("myService");

Supplier<String> decoratedSupplier =
    CircuitBreaker.decorateSupplier(circuitBreaker, () -> myService.call());

Try<String> result = Try.ofSupplier(decoratedSupplier)
    .recover(throwable -> "fallback response");
    
```

🛠️ Libraries

| Language | Library |
|----------|------------------------------------|
| Java | Resilience4j, Hystrix (deprecated) |
| .NET | Polly |
| Node.js | opossum |
| Python | pybreaker , tenacity |

The **Bulkhead pattern** is a **resilience strategy** that **isolates components** or services in a system to **prevent a failure in one part from cascading** into others.

Real-world Analogy

A **ship** is divided into watertight **bulkheads (compartments)**. If one section floods, the others stay dry – the ship remains afloat.

Similarly, in software, we create **resource boundaries** (like separate threads, processes, or pools) for different functionalities or clients.

Goal

Limit the blast radius of failures.

If one service or tenant overloads, it doesn't bring down everything else.

Types of Bulkheads

| Type | Description |
|------------------------------|---|
| Thread pool isolation | Each service/function gets a dedicated thread pool. |
| Semaphore isolation | Limits concurrent access using semaphores. |
| Service isolation | Run services on separate nodes/containers. |

Example Scenario

- Microservice A handles payments, and B handles orders.
 - Without bulkhead: high load on A → thread starvation → B also fails.
 - With bulkhead: A and B have separate thread pools → B continues to serve requests even if A is down.
-

Benefits

- **Improves fault isolation**
 - **Protects critical services**
 - **Increases system stability**
 - **Prevents cascading failures**
-

Example in Java (Resilience4j)

```

ThreadPoolBulkhead bulkhead = ThreadPoolBulkhead.ofDefaults("myService");

Supplier<CompletionStage<String>> supplier =
    ThreadPoolBulkhead.decorateSupplier(bulkhead, () -> CompletableFuture.supplyAsync
    (() -> myService.call())));

```

Combine With

- **Circuit Breaker** – block repeated failed calls.
- **Retry** – reattempt failed requests with exponential backoff.
- **Rate Limiter** – prevent abuse and overuse.

39. 

The **Retry pattern** is a **resilience technique** that automatically **retries a failed operation** (like a network call or database request) that may succeed if attempted again after a short delay.

Use Case

Retry is useful when the failure is **transient**:

- Network glitches
- Timeout errors
- Temporary service unavailability (503)
- Rate limits (e.g., HTTP 429)

How It Works

You define:

- **Max attempts** (e.g., 3 retries)
- **Delay strategy**: fixed delay, exponential backoff, or jitter

Retry Strategies

| Strategy | Description |
|----------------------------|---|
| Fixed Delay | Wait a constant time between retries |
| Exponential Backoff | Wait increases exponentially (e.g., 1s, 2s, 4s) |
| Backoff + Jitter | Adds randomness to avoid retry storms |

Example: Fixed Delay Retry (Pseudocode)

```
function retryRequest(fn, retries = 3, delay = 1000) {  
  return new Promise(async (resolve, reject) => {  
    for (let i = 0; i <= retries; i++) {  
      try {  
        const result = await fn();  
        return resolve(result);  
      } catch (error) {  
        if (i === retries) return reject(error);  
        await new Promise(r => setTimeout(r, delay));  
      }  
    }  
  });  
}
```

Good Practices

- Combine with **Circuit Breaker** to avoid retrying on known failures.
- Do **not retry** on:
 - 4xx errors (unless 429)
 - Business logic failures
- Add **timeout logic** to each retry.

Example Libraries

| Language | Library |
|----------|--|
| Java | <code>resilience4j-retry</code> |
| Node.js | <code>axios-retry</code> , <code>retry</code> |
| Python | <code>tenacity</code> , <code>retrying</code> |
| Go | <code>go-retryablehttp</code> , <code>backoff</code> |

70. Polling

Polling is a technique where the **client repeatedly sends requests** to the server at regular intervals to check if new data or changes are available.

Use Case Examples

- Checking for status updates (e.g., job status, order delivery)

- Refreshing a chat message list
- Monitoring data processing completion

How Polling Works

Client → Server → "Do you have any updates?"
(wait some time)
Client → Server → "Do you have any updates?"
...

Types of Polling

| Type | Description |
|----------------------|---|
| Short Polling | Sends a request at fixed intervals regardless of server state |
| Long Polling | Server holds the request open until data is available, then responds |

Short Polling (Basic Example)

```
setInterval(() => {
  fetch('/api/status')
    .then(res => res.json())
    .then(data => console.log(data));
}, 5000); // every 5 seconds
```

Long Polling (Pseudo Flow)

```
function poll() {
  fetch('/api/status')
    .then(res => res.json())
    .then(data => {
      console.log(data);
      poll(); // call again only after response
    });
}
poll(); // starts the loop
```

- Server keeps the request **open until an update is available**, then the client immediately starts a new one.
-

Polling vs Alternatives

| Method | When to Use |
|---------------------------------|--|
| Polling | Simple to implement, works anywhere |
| WebSockets | Realtime communication, low latency |
| Server-Sent Events (SSE) | Unidirectional push from server |
| GraphQL Subscriptions | Real-time over WebSocket in GraphQL apps |

Considerations

- Can generate **unnecessary load** on both client and server.
- Use **backoff** or **throttling** to avoid hammering the server.
- Prefer **WebSockets/SSE** when real-time is needed and supported.

71. [Load Balancing Algorithms](#)

1. Round Robin

- Cycles through servers one by one.
-  Simple and fast.
-  Doesn't account for server load.

Example:

Server A → Server B → Server C → Server A → ...

2. Weighted Round Robin

- Like Round Robin but assigns **weights** to servers based on capacity.
-  Distributes more traffic to more powerful servers.

Example:

- Server A (weight 3), B (weight 1):
A → A → A → B → A → A → A → B → ...
-

3. Least Connections

- Sends traffic to the server with the **fewest active connections**.

- Great for long-lived connections like WebSockets.
-

4. Least Response Time

- Chooses server with the **fastest response time and fewest connections**.
 - Optimizes for performance.
 - Needs real-time monitoring.
-

5. IP Hash

- Hashes the client's IP to always route to the **same server**.
 - Useful for **session persistence** (sticky sessions).
 - Doesn't rebalance well.
-

6. Random

- Randomly selects a server.
 - Simple, surprisingly effective in large systems.
-

7. Consistent Hashing

- Used in distributed systems (like **caches** or **databases**) to evenly distribute data.
 - Great for dynamic environments where nodes are added/removed often.
-

8. Resource-Based Routing

- Uses **real-time metrics** (CPU, memory, etc.) to decide where to send traffic.
 - Requires deep **monitoring + control plane**.
-

Commonly Used In:

| Algorithm | Where Used |
|----------------------|--|
| Round Robin | DNS Load Balancing, Simple Web Apps |
| Weighted Round Robin | Nginx, HAProxy |
| Least Connections | Nginx, AWS ELB, HAProxy |
| Consistent Hashing | CDNs, Distributed Caches (e.g., Memcached) |

1. [Weighted Round Robin Algorithm](#)

Weighted Round Robin (WRR) is an extension of the traditional Round Robin (RR) load balancing algorithm. The core idea of Round Robin is to distribute incoming requests evenly across all available servers in a cyclic manner. However, with **Weighted Round Robin**, servers are assigned different "weights" based on their processing power, capacity, or other metrics. This allows the load balancer to send more traffic to higher-capacity servers, ensuring that servers with greater capacity or performance characteristics handle a proportionally larger share of the load.

How it works:

1. **Assign Weights:** Each server is assigned a weight, which typically represents its processing capacity. For example, a server with a weight of 2 might handle twice as many requests as a server with a weight of 1.
2. **Distribute Requests:**
 - The load balancer cycles through the servers, but it distributes requests based on the weight assigned to each server.
 - A server with a higher weight will receive more requests per cycle than servers with a lower weight.

Example:

Imagine you have three servers with the following weights:

- **Server 1:** Weight = 3
- **Server 2:** Weight = 1
- **Server 3:** Weight = 2

Round 1:

- Server 1 gets 3 requests
- Server 2 gets 1 request
- Server 3 gets 2 requests

Round 2:

- Server 1 gets 3 requests
- Server 2 gets 1 request
- Server 3 gets 2 requests

The cycle continues, giving **Server 1** more requests than **Server 2** and **Server 3**, proportionate to its higher weight.

Key Characteristics:

1. **Weighted Allocation** : Servers with more resources (e.g., CPU, memory, or network bandwidth) can process more requests, so they get assigned higher weights.
2. **Cyclic Nature** : Like the regular round-robin algorithm, WRR ensures that requests are evenly distributed in cycles, but the distribution favors higher-weight servers.
3. **Fairness** : WRR allows more fairness in load balancing by respecting the processing capabilities of servers, unlike simple Round Robin, where all servers are treated equally regardless of their capacity.

Use Cases:

- **Web Servers** : When you have a mix of web servers with varying performance capabilities (e.g., a faster server handling more traffic than a slower one).
- **Microservices** : In a microservices architecture where services have different load capacities or resource requirements.
- **Server Pools** : When you need to efficiently distribute traffic across a pool of servers with varying CPU, memory, or bandwidth.

Algorithm Implementation:

Here's a simple way of implementing **Weighted Round Robin**:

```
class WeightedRoundRobin:
    def __init__(self, servers):
        self.servers = servers # List of tuples (server, weight)
        self.current_index = -1
        self.total_weight = sum(weight for server, weight in servers)

    def get_next_server(self):
        self.current_index = (self.current_index + 1) % self.total_weight
        current_weight_sum = 0
        for server, weight in self.servers:
            current_weight_sum += weight
            if self.current_index < current_weight_sum:
                return server
        return None

servers = [("Server 1", 3), ("Server 2", 1), ("Server 3", 2)]
wrr = WeightedRoundRobin(servers)
```

```
for _ in range(10):
    print(wrr.get_next_server())
```

Output:

```
Server 1
Server 1
Server 1
Server 2
Server 3
Server 1
Server 1
Server 1
Server 2
Server 3
```

This ensures that **Server 1** (with weight 3) gets more requests, while **Server 2** (with weight 1) gets fewer.

Summary:

- **Weighted Round Robin (WRR)** is a load balancing strategy that distributes requests to servers based on their assigned weight.
- It provides a way to fairly distribute load based on the capacity of each server, unlike standard Round Robin which assumes all servers are equal.
- It's particularly useful when you have servers with different resource capabilities and need to distribute load more intelligently.

72. [Least Privilege and Zero Trust Architecture](#)

1. Principle of Least Privilege (PoLP)

Definition:

Grant **only the minimum level of access** necessary for a user, service, or process to perform its function—**no more, no less**.

Why it's important:

- Limits the **blast radius** in case of a compromise.
- Reduces the chance of **privilege escalation** attacks.

- Protects sensitive data and systems from unintended access.
-

How to implement:

| Component | Least Privilege Example |
|------------------------|---|
| Database Access | App A can only <code>SELECT</code> from Table X, not full <code>DROP</code> |
| IAM Policies | Lambda function can only access specific S3 bucket |
| APIs | Users only get access to endpoints relevant to their role |
| Cloud Roles | Service A can't read from another service's storage |
| CI/CD | Build pipelines can deploy but not manage IAM |

2. Zero Trust Architecture (ZTA)

Definition:

"Never trust, always verify." Every access request is verified regardless of origin—**inside or outside** the network.

Core Principles:

| Principle | Description |
|---|---|
|  Verify Explicitly | Authenticate and authorize every request. |
|  Least Privilege Access | Tie in PoLP with session-based access. |
|  Assume Breach | Design as if the attacker is already in. |
|  Continuous Validation | Re-evaluate trust based on signals (IP change, device posture, behavior). |

How to implement:

| | |
|------------------------|--|
| Layer | Zero Trust Strategy |
| Network | Use VPCs, service meshes, deny-all default policies. |
| Identity | Enforce MFA, short-lived tokens, federated identity. |
| Devices | Validate device health before granting access. |
| Apps & APIs | Use mTLS, API gateways, policy engines (e.g., OPA). |
| Monitoring | Track and analyze behavior, log anomalies. |

Real-World Example: Microservice + Zero Trust + Least Privilege

User → Auth Gateway (JWT, RBAC) → Microservice A (mTLS) → DB (Encrypted)

- ✓ Gateway checks user role & issues token
- ✓ Microservice A has only read-access to DB table it owns
- ✓ All service-to-service traffic is over mTLS
- ✓ Each service has scoped, minimal permissions
- ✓ Service Mesh (e.g., Istio) enforces access policies

Tools That Help

| Category | Tools / Examples |
|----------------------|---|
| IAM | AWS IAM, GCP IAM, Azure RBAC |
| Service Mesh | Istio, Linkerd (for mTLS, traffic policies) |
| Policy Engine | OPA (Open Policy Agent), Kyverno |
| Identity | Auth0, Okta, Keycloak, Azure AD |
| Monitoring | Datadog, Prometheus, ELK, Splunk |

73. [Fallback Patterns And Fallback Strategies:](#)

1. General Failover Strategies

| Strategy | Description |
|--------------------------------------|---|
| Redis Sentinel | Provides monitoring, notification, and automatic failover to a replica. |
| Redis Cluster | Partitioned data across nodes with failover support. Good for scalability and HA. |
| Load Balancers + DNS Failover | Use smart load balancing with retry logic. |
| Persistence (AOF + RDB) | Ensure data durability; combine both for fast recovery. |
| Replication | Have one or more replicas to take over during failures. |

2. Fallback Patterns by Use Case

Use Case: Caching

Pattern: Cache-Aside (*Lazy Loading*)

```
value = redis.get(key);
if (value == null) {
```

```
    value = db.query(key);
    redis.set(key, value);
}
```

Fallback if Redis fails:

- Log the error
- Fetch from DB directly
- Optionally populate Redis when it comes back up

Redis failure doesn't block functionality — just slower.

💬 Use Case: Session Store

Fallback Pattern: Sticky Fallback or Local Session

- In the event of Redis outage:
 - Store session in app memory (only for short periods)
 - Notify user to re-authenticate with grace
 - Load from a secondary store like DB

High Availability Strategy:

Use **Redis Sentinel** to promote a replica instantly.

🚦 Use Case: Rate Limiting / Quota Tracking

Pattern: Fail Open or Fail Closed

- **Fail Open:** Let requests pass if Redis is down.
- **Fail Closed:** Block requests cautiously, safer for security-sensitive APIs.

```
try {
    allowed = redis.eval(luaScript);
} catch (RedisDownException e) {
    log.warn("Redis down, allowing request");
    allowed = true; // fail open
}
```

You can also:

- Fall back to in-memory counters (not ideal for distributed systems).
- Use **token bucket** in app memory temporarily.

Use Case: Pub/Sub Messaging

Fallback Strategies:

- Use **Redis Streams** instead of basic Pub/Sub (they're persistent).
 - Or better: shift critical messaging to **Kafka / RabbitMQ**.
 - If Redis fails:
 - Buffer messages temporarily (disk or memory)
 - Retry once Redis is healthy again
-

Handling Redis Failures in Code

- **Retry with exponential backoff**
 - **Use a circuit breaker:**
 - Open after X failures
 - Wait for cool-off time before retrying
 - Example libraries: Resilience4j, [Hystrix](#)
-

Bonus: Readiness & Liveness Probes

- Use Redis `PING` for health checks.
- Kubernetes probes can remove unhealthy Redis pods automatically.
- Monitor metrics like:
 - `connected_clients`
 - `used_memory`
 - `rdb_last_save_time`

75. What if redis goes down in the system?

If **Redis** goes down in a system, the **impact depends on how your system uses Redis**. Redis is often used as a cache, pub/sub system, rate limiter, session store, or even as a primary datastore in some cases.

Here's a breakdown of **what happens and what you can do**:

What Could Go Wrong?

1. As a Cache (e.g., for DB queries)

- **Impact:** Slower response times — system falls back to DB.
 - **Severity:** Low to Moderate
 - **Mitigation:**
 - Use **graceful fallbacks** to DB (cache-aside pattern).
 - Add **circuit breakers** to avoid overwhelming the DB.
-

2. As a Session Store (e.g., login sessions)

- **Impact:** Users might get logged out or see errors.
 - **Severity:** Moderate
 - **Mitigation:**
 - Use **sticky sessions**, or fallback to local session storage.
 - Persist sessions in a **high-availability Redis setup**.
-

3. As a Message Broker (Pub/Sub)

- **Impact:** Message loss or system components stop communicating.
 - **Severity:** High
 - **Mitigation:**
 - Use **Redis Streams** (persistence) instead of ephemeral Pub/Sub.
 - Consider **Kafka or RabbitMQ** for critical systems.
-

4. As a Rate Limiter

- **Impact:** Either unthrottled or blocked traffic.
 - **Severity:** High
 - **Mitigation:**
 - Fail open with logging or degrade gracefully.
 - Fallback to **in-memory counters** temporarily.
-

5. As a Primary Store

- **Impact:**  Data loss (if not persisted), system crash.
- **Severity:** Very High

- **Mitigation:**
 - Enable **RDB/AOF persistence** in Redis.
 - Use **Redis Sentinel** or **Redis Cluster** for high availability.
 - Backup regularly and failover smartly.
-

General Best Practices

| Strategy | Description |
|---|---|
| Redis Sentinel / Cluster | Auto failover, master election, replication |
| Persistence (AOF, RDB) | Durable storage to restore data |
| Health checks & alerts | Proactive monitoring of Redis |
| Graceful fallbacks | Don't crash if Redis is unavailable |
| Connection pooling & retry logic | Handle transient failures smoothly |
| Circuit Breakers | Avoid hammering Redis repeatedly if it's down |

76. gRPC Vs REST

| Feature | gRPC | REST |
|------------------|--|------------------------------------|
| Protocol | HTTP/2 | HTTP/1.1 |
| Message Format | Protocol Buffers (binary) | JSON (text) |
| Communication | Unary, Streaming (client/server/bidirectional) | Request-response |
| Performance | Faster (compact, binary) | Slower (text-based, more overhead) |
| Schema | <code>.proto</code> file (strongly typed) | No enforced schema |
| Language Support | Multi-language (via code-gen) | Multi-language (manually handled) |
| Browser Support | Not directly (needs proxy/adapter) | Native |
| Tooling | Great for internal microservices | Great for public APIs |

Internal Implementation

gRPC

- Built on **HTTP/2**, enabling:
 - **Multiplexing** (multiple requests/responses over same TCP)
 - **Header compression**
 - **Bi-directional streaming**
- Uses **Protocol Buffers (protobuf)** for defining service interfaces and messages:

```

service Greeter {
    rpc SayHello (HelloRequest) returns (HelloReply);
}

message HelloRequest {
    string name = 1;
}

message HelloReply {
    string message = 1;
}

```

- Server generates **stub** and **skeleton** code:
 - Client uses the stub like a local method call.
 - Server implements the service interface.
- Supports:
 - **Unary RPC**
 - **Client Streaming**
 - **Server Streaming**
 - **Bidirectional Streaming**

◆ REST

- Based on **HTTP 1.1**, stateless.
- Uses **URLs for resources** (nouns) and **HTTP verbs** (GET, POST, PUT, DELETE).
- Payloads are in **JSON**, optionally XML or plain text.
- Human-readable but more **overhead** due to:
 - Larger payload sizes
 - Lack of strict typing
- Example:

```

GET /users/123
POST /users

```

⚡ Performance Comparison

| | | |
|-------------------|----------------------|-------------------------------|
| Feature | gRPC | REST |
| Speed | Faster | Slower |
| Payload Size | Small (binary) | Large (text) |
| Streaming Support | Excellent (built-in) | Needs long-polling/websockets |
| Compression | Built-in via HTTP/2 | Optional (e.g. gzip) |

When to Use What?

| Use Case | Recommendation |
|---------------------------------|-------------------------------------|
| Internal microservices | gRPC (performance, contract) |
| Public APIs for web clients | REST (browser native) |
| Need for streaming | gRPC |
| Human-readability/debuggability | REST |

Compression and decompression of binary content generated using protobuf and IDL makes it faster

⚙️ 1. REST (Express)

server.js (REST)

```
js Copy Edit  
  
const express = require('express');  
const app = express();  
app.use(express.json());  
  
app.post('/sayHello', (req, res) => {  
  const { name } = req.body;  
  res.send({ message: `Hello, ${name}!` });  
});  
  
app.listen(3000, () => console.log('REST server running on port 3000'));
```

client.js (REST)

```
js Copy Edit  
  
const axios = require('axios');  
  
(async () => {  
  const response = await axios.post('http://localhost:3000/sayHello', {  
    name: 'Sai'  
  });  
  console.log(response.data); // { message: "Hello, Sai!" }  
})();
```

2. gRPC

greeter.proto

```
proto
syntax = "proto3";

service Greeter {
    rpc SayHello (HelloRequest) returns (HelloReply);
}

message HelloRequest {
    string name = 1;
}

message HelloReply {
    string message = 1;
}
```

Copy Edit

server.js (gRPC)

```
js
const grpc = require('@grpc/grpc-js');
const protoLoader = require('@grpc/proto-loader');

const packageDefinition = protoLoader.loadSync('greeter.proto');
const proto = grpc.loadPackageDefinition(packageDefinition);

function sayHello(call, callback) {
    callback(null, { message: `Hello, ${call.request.name}!` });
}
```

Copy Edit

```

const server = new grpc.Server();
server.addService(proto.Greeter.service, { SayHello: sayHello });
server.bindAsync('127.0.0.1:50051', grpc.ServerCredentials.createInsecure(), () => {
  server.start();
  console.log('gRPC server running on port 50051');
});

```

client.js (gRPC)

js

Copy Edit

```

const grpc = require('@grpc/grpc-js');
const protoLoader = require('@grpc/proto-loader');

const packageDefinition = protoLoader.loadSync('greeter.proto');
const proto = grpc.loadPackageDefinition(packageDefinition);

const client = new proto.Greeter('localhost:50051', grpc.credentials.createInsecure());

client.SayHello({ name: 'Sai' }, (err, response) => {
  console.log(response); // { message: 'Hello, Sai!' }
});

```

🚀 Summary

| Feature | REST (Express) | gRPC |
|-------------|--------------------------------------|------------------------------|
| Simplicity | ✓ Easy | ✗ Requires .proto & setup |
| Performance | 🐢 Slower (JSON + HTTP 1.1) | ⚡ Faster (ProtoBuf + HTTP/2) |
| Type Safety | ✗ Loose typing | ✓ Strongly typed |
| Streaming | ✗ Needs custom WebSockets or polling | ✓ Built-in support |

76. Concurrency Patterns

🔄 1. Worker Pool / Thread Pool

- **Use:** Limit the number of concurrent tasks to avoid resource exhaustion.
- **How it works:** A fixed set of worker threads pick tasks from a queue and process them.

- **Example:** Node.js `cluster`, Java `ExecutorService`, Python `ThreadPoolExecutor`.
-

2. Producer–Consumer

- **Use:** Decouples task production and processing.
 - **How it works:** One or more producers put tasks into a shared queue; one or more consumers process them.
 - **Real-world analogy:** Restaurant kitchen (orders vs chefs).
 - **Tools:** BlockingQueue (Java), Channel (Golang), RabbitMQ/Kafka.
-

3. Fan-Out / Fan-In

- **Fan-Out:** Split a task into multiple subtasks to be done in parallel.
 - **Fan-In:** Combine the results of parallel subtasks.
 - **Use:** Parallel map/reduce, distributed jobs.
 - **Example:** Serverless fan-out via AWS Lambda → results aggregated later.
-

4. Backpressure

- **Use:** Prevent producers from overwhelming consumers.
 - **How it works:** Slow down or block the producer if the queue is full.
 - **Implemented in:** RxJS, Akka Streams, Kafka, Node Streams.
-

5. Event Loop (Reactor Pattern)

- **Use:** Handle many I/O-bound tasks with a single thread (non-blocking).
 - **Common in:** Node.js, JavaScript browser runtime, Python asyncio.
 - **Core idea:** Events are handled asynchronously via callbacks or promises.
-

6. Semaphore / Mutex (Mutual Exclusion)

- **Use:** Prevent race conditions in shared resources.
 - **Semaphore:** Limits the number of threads accessing a resource.
 - **Mutex:** Allows only one thread at a time.
-

7. Fork–Join

- **Use:** Parallelism in CPU-intensive tasks.
-

- **How it works:** Split tasks into smaller subtasks (fork), solve them in parallel, then combine results (join).
 - **Example:** Java ForkJoinPool, parallel loops in Python, C++.
-

8. Scheduler Pattern

- **Use:** Run tasks at a fixed time or interval.
 - **Examples:** Cron jobs, retry tasks, cleanup jobs.
 - **Tools:** Quartz Scheduler (Java), Node `node-cron`, Celery beat (Python).
-



9. Barrier / Latch

- **Barrier:** Wait for a group of threads to reach a common point before proceeding.
 - **Latch:** Wait for one or more threads to complete before continuing.
 - **Used in:** Coordinated starts, phased execution.
-

10. Actor Model

- **Use:** Concurrency through message passing instead of shared memory.
- **Each actor:** Has state + mailbox; handles one message at a time.
- **Used in:** Akka (Scala/Java), Erlang, Orleans (.NET), Actix (Rust).